

# Selendra Neural Network (SNN) Technical Whitepaper

## Introduction

This document provides a detailed technical blueprint for the implementation of the Selendra Neural Network (SNN), translating its conceptual design into actionable engineering specifications. It addresses key aspects of implementability, including algorithms, data structures, protocols, cryptographic primitives, component breakdown, challenges, security, parameterization, integration, research needs, and comparisons with existing systems for each major component of the SNN. The SNN aims to deliver a highly scalable, interoperable, user-friendly, and developer-centric blockchain platform by leveraging a novel hybrid DAG-based consensus mechanism, dynamic sharding, native account abstraction, and a robust economic model.

---

## 1. Core Ledger Structure & Consensus (DAG-centric Exploration)

This section details the implementation plan for SNN's core ledger, focusing on its unique hybrid DAG model, transaction processing, and consensus mechanisms.

### 1.1. DAG Model Selection/Innovation: The Synaptic Ledger

The Synaptic Ledger is envisioned as a hybrid Proof-of-Stake (PoS) Directed Acyclic Graph (DAG) that draws inspiration from Avalanche, Tangle, and Block-Lattice architectures. This design aims to achieve high throughput, low latency, and robust security.

#### 1.1.1. Proposed Hybrid PoS DAG Structure

- **Specific Algorithms & Data Structures:**
  - **Data Structure (DAG):** The primary data structure will be a DAG where vertices represent "Transaction Units" (TUs). A TU can be a single transaction or a micro-block containing multiple transactions. Each TU will reference  $k$  parent TUs (e.g.,  $k=2$  to  $5$ ).
    - *Block-Lattice Influence:* To enhance parallelism and user-level ordering, each account will maintain its own chain of TUs ("account-

synapse”). TUs within an account-synapse are causally ordered. Cross-account interactions are represented by TUs that reference parents from different account-synapses.

- **Graph Representation:** Adjacency lists stored in a key-value database (e.g., RocksDB, LevelDB) optimized for graph traversals. Each TU will store: TU ID (hash), parent IDs, payload (transactions), issuer’s signature, timestamp, and consensus metadata.
- **Data Structure (Validator Sets):** A dynamically updated list of registered validators, their stake amounts, and public keys, managed by a staking contract/module. Stored on-chain and replicated.
- **Algorithm (PoS Mechanism):** Bonded Proof-of-Stake. Validators lock NEURON tokens to participate. The probability of being selected to issue a TU or participate in consensus sampling rounds is proportional to their stake. Slashing conditions will apply for malicious behavior.
- **Algorithm (Tip Selection):** A weighted random walk algorithm, influenced by Tangle’s Unweighted Random Walk, but modified to incorporate validator stake/reputation and TU “health” (e.g., TUs with higher cumulative weight, or those referenced by reputable validators). This aims to ensure the DAG grows securely and efficiently, preventing lazy tip selection and promoting confirmation of older TUs.
- **Protocol Details & Message Flows:**
  - **TU Creation & Propagation:**
    1. User/Client creates a transaction.
    2. Transaction is submitted to a validator node.
    3. Validator validates the transaction.
    4. Validator selects  $k$  parent TUs using the tip selection algorithm.
    5. Validator bundles the transaction(s) into a TU, signs it, and adds it to its local DAG view.
    6. Validator gossips the new TU to its connected peers using a P2P protocol (e.g., libp2p with gossipsub).
  - **DAG Synchronization:** Nodes continuously exchange TU information with peers to maintain a consistent view of the DAG.
- **Cryptographic Primitives:**
  - **Hash Functions:** BLAKE3 or SHA3-256 for TU IDs, Merkle tree roots.
  - **Signature Schemes:** EdDSA (Ed25519) for signing TUs. Schnorr signatures for potential aggregation.
- **Component Breakdown:**
  - **P2P Networking Layer:** (e.g., using libp2p).
  - **DAG Storage Engine:** (e.g., built on RocksDB).
  - **Transaction Pool (Mempool).**

- **Staking Management Module.**
- **TU Construction & Validation Module.**
- **Implementation Challenges & Trade-offs:**
  - **Challenge (Tip Selection):** Designing a manipulation-resistant tip selection algorithm.
  - **Challenge (Data Storage):** Efficiently storing and querying the DAG.
  - **Trade-off (TU Size):** Single transaction TUs vs. micro-blocks.
  - **Trade-off (Account-Synapse Complexity):** Parallelism vs. complexity of cross-account interactions.
- **Security Considerations (Implementation Specific):**
  - **Long-Range Attacks:** Mitigation: Checkpointing, stake lock-up periods.
  - **Selfish Mining/Issuance:** Mitigation: Robust tip selection, network monitoring.
  - **DAG Pruning Complexity:** Ensuring safe pruning.
- **Parameterization & Configuration:**
  - $k$  (parents per TU).
  - Validator stake thresholds.
  - Epoch duration for validator rotation.
  - Tip selection parameters.
  - Gossip protocol parameters.
- **Integration Points:**
  - **Consensus Engine.**
  - **Transaction Execution Layer.**
  - **Governance Module.**
- **Research & Prototyping Needs:**
  - Simulation of the hybrid DAG model and tip selection.
  - Benchmarking database solutions for DAG storage.
  - Prototyping account-synapse interaction logic.
- **Comparison with Existing Implementations:**
  - **Fantom (Opera), Hedera Hashgraph, Aleph Zero:** SNN's innovation lies in the specific synthesis of block-lattice, Tangle, and Avalanche influences, aiming for enhanced parallelism and a unique PoS integration.

### 1.1.2. Transaction Ordering (Causal, Consensus Timestamping & Epochs)

- **Specific Algorithms & Data Structures:**
  - **Data Structure (TU):** TUs contain validator timestamps and parent references.

- **Algorithm (Causal Ordering):** Defined by DAG parent-child relationships.
- **Algorithm (Consensus Timestamping):** Weighted median of timestamps proposed by validators in a consensus round for a TU or an epoch batch.
- **Data Structure (Epochs):** Timeline divided into epochs for synchronization, validator rotation, and batch timestamping. Epoch marker TUs.
- **Protocol Details & Message Flows:**
  - **Timestamping:** Validator includes local timestamp in TU. Consensus process agrees on a final consensus timestamp (e.g., median from final sampling round or of TUs in confirmation set).
  - **Epoch Transition:** Triggered by committee or smart contract, potentially issuing special epoch boundary TUs.
- **Cryptographic Primitives:**
  - Standard hashes and signatures. Threshold signatures for committee-generated epoch markers.
- **Component Breakdown:**
  - **Time Synchronization Service (e.g., NTP).**
  - **Consensus Module (for timestamp agreement).**
  - **Epoch Management Module.**
- **Implementation Challenges & Trade-offs:**
  - **Challenge (Timestamp Manipulation):** Median-based consensus mitigates but doesn't eliminate.
  - **Trade-off (Epoch Length):** Shorter epochs (faster rotation, granular timestamps) vs. higher overhead.
  - **Trade-off (Ordering Granularity):** Causal order (simple, incomplete) vs. Consensus timestamping (total order, complex).
- **Security Considerations (Implementation Specific):**
  - **Timestamp Attacks:** Resilience against skewed clocks.
  - **Epoch Boundary Manipulation:** Robust handling to prevent stalling.
- **Parameterization & Configuration:**
  - Epoch duration.
  - Timestamp consensus threshold (e.g., 2/3 stake).
  - Max allowable clock skew.
- **Integration Points:**
  - **Execution Environment (relies on final order).**
  - **Consensus Engine (provides timestamp agreement).**
  - **Slashing Module (penalize outlier timestamps).**

- **Research & Prototyping Needs:**
  - MEV mitigation for DAGs with consensus timestamping.
  - Simulation of epoch lengths and timestamping strategies.
- **Comparison with Existing Implementations:**
  - **Hedera Hashgraph, Solana (PoH):** SNN uses consensus *on* timestamps within epochs, differing from Hedera's continuous model and Solana's pre-consensus clock.
  - **SNN Innovation:** Balancing DAG-native causal order with epoch-based consensus timestamping for speed, fairness, and strong ordering.

### 1.1.3. Consensus (Finality via Avalanche-like sampling)

- **Specific Algorithms & Data Structures:**
  - **Algorithm (Consensus):** Snowball/Avalanche protocol. Repeated random, stake-weighted sampling of validators to determine preference for TUs.
  - **Data Structure (Confidence):** Per-validator map of TU IDs to confidence counters and preference status.
  - **Algorithm (Validator Sampling):** Secure random selection (potentially VRF-based).
- **Protocol Details & Message Flows:**
  - **Query Round:** Validator A queries  $k$  peers on TU X. Peers respond with preference. A updates its preference/confidence based on  $\alpha * k$  majority. Process repeats.
  - **Finality:** Probabilistic finality when a TU is accepted with high confidence by a validator. Network-wide finality can be achieved when a large stake portion accepts it, possibly observed by a finality gadget.
- **Cryptographic Primitives:**
  - Standard signatures for query/response integrity.
  - VRF for fair validator sampling (recommended).
- **Component Breakdown:**
  - **Consensus Engine (Avalanche logic).**
  - **Networking Module (for query/response).**
  - **Finality Gadget (Optional, for stronger finality).**
- **Implementation Challenges & Trade-offs:**
  - **Challenge (Parameter Tuning):**  $k$  (sample size),  $\alpha$  (quorum),  $\beta$  (confidence) are critical for safety/liveness.
  - **Challenge (Network Latency):** Performance sensitive to communication rounds.

- **Trade-off (Probabilistic vs. Absolute Finality):** Avalanche offers fast probabilistic finality. Absolute finality adds complexity.
- **Security Considerations (Implementation Specific):**
  - **Sybil Attacks (Sampling):** Mitigated by stake-weighted sampling.
  - **Colluding Minority:** Can stall or conflict if parameters are weak.
  - **Liveness under Attack:** Avalanche is generally robust.
- **Parameterization & Configuration:**
  - `k` , `alpha` , `beta1` , `beta2` (Avalanche parameters).
  - Number of consensus rounds.
  - Concurrent consensus instances per validator.
- **Integration Points:**
  - **DAG Storage/Transaction Processing.**
  - **Execution Environment.**
  - **Slashing Module (for non-participation/malice in consensus).**
- **Research & Prototyping Needs:**
  - Simulation for optimal Avalanche parameter tuning.
  - Investigating lightweight finality gadget integration.
- **Comparison with Existing Implementations:**
  - **Avalanche Platform (AVAX):** SNN directly uses Avalanche principles, differentiating via the underlying Synaptic Ledger and integration with other SNN features.
  - **SNN Innovation:** Applying Avalanche consensus to its unique hybrid DAG structure.

#### 1.1.4. Security (Double-Spends, Sybil Attacks, Parasitic Chains, MEV Resistance - Threshold Encryption, Fair Ordering)

- **Specific Algorithms & Data Structures:**
  - **Double-Spends:** Resolved by DAG structure and Avalanche consensus finalizing only one conflicting TU. Account state/UTXO set used for detection.
  - **Sybil Attacks:** Mitigated by PoS and staking registry.
  - **Parasitic Chains / Lazy Tips:** Mitigated by enhanced tip selection (see 1.1.1) and Avalanche consensus requiring broad acceptance.
  - **MEV Resistance (Threshold Encryption):**
    - *Algorithm:* Threshold Encryption Scheme (TES) (e.g., based on BLS or IBE). Transactions encrypted with a common public key, decrypted by a committee post-ordering and finalization. Distributed Key Generation (DKG) for committee keys.
    - *Data Structure:* Encrypted transaction pool, decryption key shares.

- **MEV Resistance (Fair Ordering):**
  - *Algorithm:* Consensus-driven ordering based on consensus timestamps. Deterministic tie-breaking (e.g., TU hash) for identical timestamps. Alternatively, commit-reveal schemes or VDF-based approaches for proposers.
- **Protocol Details & Message Flows:**
  - **Double-Spend Resolution:** Conflicting TUs enter consensus; one is finalized, others rejected.
  - **Threshold Encrypted Transaction Submission:** User encrypts tx -> submits -> encrypted tx ordered by consensus -> decryption committee decrypts -> decrypted tx executed.
- **Cryptographic Primitives:**
  - **Threshold Encryption:** BLS-based TES or similar, DKG protocols.
  - **Standard Primitives:** EdDSA, BLAKE3.
  - **VRFs/VDFs (Optional for advanced Fair Ordering).**
- **Component Breakdown:**
  - **Conflict Detection Module.**
  - **Threshold Cryptography Module (DKG, encryption, decryption).**
  - **Fair Ordering Service/Logic.**
  - **Slashing Module (for MEV misbehavior, premature decryption).**
- **Implementation Challenges & Trade-offs:**
  - **Challenge (Threshold Encryption):** DKG complexity, performance overhead, secure committee management.
  - **Challenge (Fair Ordering):** Defining and achieving true fairness is hard and can impact performance.
  - **Trade-off (MEV Resistance vs. Complexity/Performance):** Strong MEV resistance adds significant overhead.
- **Security Considerations (Implementation Specific):**
  - **Threshold Encryption Collusion:**  $t+1$  collusion in decryption committee.
  - **Side Channels (Encrypted Transactions):** Leakage from size/metadata.
  - **Fair Ordering Bypass:** Potential for sophisticated attacks.
- **Parameterization & Configuration:**
  - Threshold  $t$  and  $n$  for decryption committee.
  - Threshold key refresh period.
  - Fair ordering parameters (e.g., batching windows).
- **Integration Points:**
  - **Transaction Pool (handles encrypted/plain tx).**
  - **Consensus Engine (ordering is input to decryption).**

- **Execution Environment (executes post-decryption).**
- **Governance (manages decryption committee).**
- **Research & Prototyping Needs:**
  - Prototyping TES and DKG for performance/security assessment.
  - Simulation of fair ordering algorithms for MEV impact.
  - Monitoring tools for parasitic chains/MEV.
- **Comparison with Existing Implementations:**
  - **Secret Network (TEEs), Shutter Network, Ferveo (Threshold Crypto for MEV):** SNN draws from TES approaches, integrating them into a DAG.
  - **Chainlink FSS, Flashbots:** SNN aims to neutralize MEV via encryption/fair ordering, differing from Flashbots' auction model.
  - **SNN Innovation:** Comprehensive base-layer MEV resistance by combining PoS DAG with threshold encryption and consensus-driven fair ordering.

## 1.2. Scalability Advantages & Challenges

This section focuses on how the implementation choices for the Synaptic Ledger (Section 1.1) directly impact SNN's scalability, and the challenges that arise.

- **Specific Algorithms & Data Structures (Impacting Scalability):**
  - **DAG Structure:** The ability to process TUs in parallel, without waiting for a global block, is the primary scalability driver. The “account-synapse” concept (1.1.1) further enhances this by allowing independent transaction processing for different accounts until an interaction occurs.
  - **Avalanche Consensus:** Its sampling nature means not all validators need to process every TU for consensus, reducing communication overhead compared to classical BFT.
  - **Tip Selection Algorithm:** An efficient tip selection algorithm ensures the DAG remains healthy and TUs are confirmed quickly, contributing to high throughput.
  - **Data Structures for DAG Storage:** Using optimized key-value stores (e.g., RocksDB with custom indexing for parent/child lookups, tip traversal) is crucial for handling a large, growing DAG.
- **Protocol Details & Message Flows (Impacting Scalability):**
  - **Gossip Protocol (TU Propagation):** Efficient gossip (e.g., libp2p's gossipsub with topic-based subscriptions, potentially sharded by DSCs later) is vital to quickly disseminate TUs without overwhelming nodes.
  - **Consensus Message Flow:** Avalanche's query mechanism is lightweight. The number of rounds and sample size (  $k$  ) directly impact latency and



message load.

- **Component Breakdown (Impacting Scalability):**
  - **P2P Networking Layer:** Must handle high message rates and efficient routing.
  - **DAG Storage Engine:** Read/write performance is a bottleneck.
  - **Execution Environment:** If TUs contain complex smart contracts, execution speed per TU becomes a factor. (More in Section 5).
- **Implementation Challenges & Trade-offs (Scalability Specific):**
  - **Challenge (State Bloat):** A highly active DAG can lead to rapid state growth. Addressed by State Management (2.3).
  - **Challenge (Bandwidth Consumption):** Gossiping all TUs to all nodes in a very high-throughput scenario is unsustainable. This motivates DSCs (2.2).
  - **Challenge (Processing Bottlenecks):** Even with parallel TU processing, individual nodes have CPU, I/O, and memory limits. Node optimization and potentially heterogeneous node roles are important.
  - **Challenge (Parent Selection Under Load):** Ensuring TUs select good parents quickly when the DAG is very wide and growing fast.
  - **Trade-off (Confirmation Latency vs. Throughput):** Aggressive parameter tuning for Avalanche might increase throughput but could slightly increase the time to high confidence for individual TUs or reduce resilience.
  - **Trade-off (Redundancy vs. Efficiency):** How many peers a TU is gossiped to, or how many validators store the full DAG, impacts resilience vs. storage/bandwidth costs.
- **Security Considerations (Scalability Specific):**
  - **Increased Attack Surface (DoS):** Higher throughput means more potential vectors for spamming the network with TUs that consume resources but have no economic value. Fee mechanisms (4.2) and PoS are primary mitigations.
  - **Amplification Attacks:** Poorly designed gossip or query protocols could be exploited.
- **Parameterization & Configuration (Scalability Specific):**
  - Gossip protocol fanout, message quotas.
  - Avalanche parameters (  $k$  ,  $\alpha$  ,  $\beta$  ).
  - Database cache sizes, compaction strategies.
  - Limits on TU size or computational complexity.
- **Integration Points:**
  - **Dynamic Synaptic Clusters (DSCs) (2.2):** Essential for horizontal scaling beyond single DAG limits.

- **State Management (2.3):** Crucial for handling the state growth from high throughput.
- **Fee Model (4.2):** To prevent spam and align resource usage with cost.
- **Research & Prototyping Needs:**
  - Extensive benchmarking of node software under high load.
  - Simulation of DAG growth and health under various transaction patterns and network conditions.
  - Developing adaptive gossip protocols that respond to network load.
- **Comparison with Existing Implementations:**
  - **Solana:** Achieves high TPS via PoH, pipelined processing, and a focus on optimized single-leader block production. SNN's DAG approach is leaderless, aiming for similar throughput via parallel TU processing and Avalanche.
  - **Aptos/Sui:** Use parallel execution with optimistic concurrency control. SNN's DAG inherently allows parallel validation/consensus up to a point, with execution being the next step. The "account-synapse" has similarities to object-centric models in terms of enabling parallelization for non-contentious operations.
  - **SNN Innovation:** The specific hybrid DAG model combined with Avalanche, aiming for a balance of high throughput, low latency, decentralization, and security without a single leader.

### 1.3. Transaction Processing

This covers the lifecycle of a transaction from submission to final confirmation within the Synaptic Ledger.

- **Specific Algorithms & Data Structures:**
  - **Data Structure (Transaction):** Standard fields: sender, receiver, value, data (payload for smart contracts), nonce/sequence number (per account-synapse), signature, fee offer.
  - **Data Structure (Transaction Unit - TU):** As defined in 1.1.1. Can contain one or more transactions.
  - **Data Structure (Mempool/Transaction Pool):** Per-validator, priority queue for pending transactions, potentially sharded or organized by fee rate and account-synapse.
  - **Algorithm (Transaction Validation):**
    1. Syntactic Validity: Correct format, fields present.
    2. Signature Verification: Using sender's public key.
    3. Semantic Validity: Sufficient balance for value + fee, correct nonce, sender permissions. For smart contracts, basic checks before adding to TU.

- **Algorithm (TU Selection for Consensus):** Validators select TUs from their local DAG view to participate in Avalanche sampling rounds. Preference is given to TUs that extend the “heaviest” or most “virtuous” paths.
- **Protocol Details & Message Flows:**
  - **1. Submission:**
    - User crafts a transaction, signs it.
    - User submits it to one or more SNN validator nodes via an RPC interface.
  - **2. Propagation (Initial):**
    - Receiving validator validates the transaction. If valid, adds to its mempool.
    - Validator may gossip the raw transaction to a few peers, especially if it doesn’t plan to issue a TU immediately.
  - **3. Validation & DAG Inclusion (TU Creation):**
    - Validator selects transactions from its mempool (e.g., highest fee first, respecting account-synapse order).
    - Validator selects `k` parent TUs from its DAG view (tip selection - 1.1.1).
    - Validator creates a new TU, bundling the selected transactions, adds parent references, signs it.
    - Validator adds the TU to its local DAG and gossips it to peers (1.1.1).
  - **4. Propagation (TU Gossip):**
    - Peers receiving the TU validate it (parents exist, issuer is valid, transactions within are plausible). If valid, add to their local DAG and re-gossip.
  - **5. Consensus & Confirmation (Avalanche Sampling):**
    - Validators run Avalanche sampling rounds on TUs in their view (1.1.3).
    - A TU becomes “preferred” by a validator once its confidence counter ( `beta1` ) is met.
    - A TU becomes “accepted” (locally confirmed) once its consecutive success counter ( `beta2` ) is met.
    - Transactions within an accepted TU are considered confirmed by that validator. Network-wide finality is probabilistic and strengthens as more stake accepts the TU.
  - **6. Execution:** Once a TU is confirmed and ordered (1.1.2), its transactions are executed against the state. (More in Section 5).
- **Cryptographic Primitives:**
  - EdDSA for transaction signatures.

- BLAKE3/SHA3-256 for hashing transactions and TUs.
- **Component Breakdown:**
  - **RPC Interface:** For transaction submission.
  - **Mempool Management Module.**
  - **Transaction Validation Module.**
  - **TU Construction Module.**
  - **DAG Update & Gossip Module.**
  - **Consensus Engine (Avalanche).**
  - **Execution Environment Interface.**
- **Implementation Challenges & Trade-offs:**
  - **Challenge (Mempool Management):** Preventing spam, efficient prioritization, especially with account-synapses.
  - **Challenge (Race Conditions in TU Creation):** Multiple validators might try to include the same transaction from their mempools. The first TU to get confirmed wins; others are discarded or their transactions re-queued if still valid.
  - **Trade-off (TU Size vs. Latency):** Larger TUs (more txns) can increase throughput but might slightly increase propagation and validation time for that TU.
  - **Trade-off (Validation Depth before Gossiping TU):** Quick gossip vs. ensuring TUs are mostly valid to save network resources.
- **Security Considerations (Implementation Specific):**
  - **Transaction Spamming at RPC/Mempool:** Rate limiting, fee requirements.
  - **Replay Attacks:** Nonce/sequence number per account-synapse prevents this.
  - **Invalid TU Propagation:** Robust validation by receiving peers. Slashing for validators issuing provably invalid TUs.
- **Parameterization & Configuration:**
  - Mempool size limits (count, bytes).
  - Min/max TU size (transaction count, bytes).
  - Transaction fee parameters (see 4.2).
  - RPC request limits.
- **Integration Points:**
  - **All other modules in Section 1.**
  - **Execution Environment (Section 5).**
  - **Fee Market (Section 4.2).**
  - **P2P Network Layer.**

- **Research & Prototyping Needs:**
  - Optimized mempool designs for DAGs with account-based parallelism.
  - Strategies for efficient TU batching by validators.
- **Comparison with Existing Implementations:**
  - **Bitcoin/Ethereum:** Similar submission to mempool, but SNN validators proactively pull transactions into TUs for DAG inclusion rather than miners/block proposers selecting for a single block.
  - **DAG Systems (e.g., Fantom, Hedera):** Broadly similar flow. SNN's specifics are in the hybrid DAG structure, Avalanche consensus, and account-synapse model influencing validation and inclusion.
  - **SNN Innovation:** The tight coupling of transaction processing with the Synaptic Ledger's specific DAG structure and Avalanche consensus, aiming for rapid inclusion and confirmation.

## 1.4. Alternative Considerations

Exploring fallback or future enhancements for the core ledger structure.

- **Micro-Blocks/DAG-Chains:**
  - **Concept:** TUs are essentially micro-blocks. A “DAG-Chain” could imply a more structured approach where sequences of TUs from specific, highly active validators or applications form mini-chains that are then cross-referenced in the wider DAG. This is somewhat captured by the “account-synapse” idea but could be formalized further for specific use cases (e.g., a high-frequency trading DSC).
  - **When/How to Implement:**
    - *When:* If certain applications generate a very high volume of related transactions that could benefit from localized, faster sequencing before broader DAG consensus, or if the general DAG becomes too complex to efficiently order for all use cases.
    - *How:* Define specific TU types or flags that indicate they belong to a particular DAG-Chain. Validators participating in that DAG-Chain would prioritize linking these TUs. Broader SNN consensus would still finalize these chains. This could be an overlay structure or a specialized type of DSC.
  - **Specific Algorithms & Data Structures:** Each DAG-Chain might have its own local leader rotation or sequencing rule, with its root hash or latest TU anchored into the main Synaptic Ledger.
  - **Challenges:** Increased complexity, potential for centralization within a DAG-Chain if not designed carefully.
- **Hybrid DAG + Execution Layer (Rollup-like internal sharding):**
  - **Concept:** The main Synaptic Ledger (L1 DAG) focuses primarily on ordering and security for batches of transactions or state commitments

from multiple, parallel “Execution Layers” (ELs). These ELs could be specialized shards, application-specific chains, or even different VM environments that process transactions independently and then submit proofs/state roots to the L1 DAG.

- **When/How to Implement:**
  - *When:* If the core SNN execution environment (WASM on DAG, Section 5.1) itself becomes a bottleneck despite DAG parallelism, or if diverse execution environments are highly desired. This is a natural evolution towards a more explicitly sharded or rollup-centric architecture, using the DAG as a data availability and settlement layer.
  - *How:*
    1. Define interfaces for ELs to submit transaction batches/state roots to the SNN L1 DAG.
    2. The L1 DAG validators would primarily validate these EL commitments and order them via Avalanche consensus.
    3. Fraud proofs or validity proofs (ZK-SNARKs) from ELs would be verified by L1 DAG validators or specialized verifier nodes.
    4. Cross-EL communication would occur via the L1 DAG.
- **Specific Algorithms & Data Structures:** Merkle trees for EL state commitments, ZK circuits for validity proofs, or challenge-response protocols for fraud proofs.
- **Challenges:** Significant architectural complexity, ensuring secure and efficient proof systems, managing cross-EL data consistency.
- **Research & Prototyping Needs for Alternatives:**
  - Simulating performance and complexity of DAG-Chains for specific use cases.
  - Developing proof-of-concept for an EL submitting ZK proofs of its state transitions to a simplified SNN L1 DAG. This aligns with the ZK-Rollup bridge concept (3.1.3) but for internal scaling.
- **Integration Points:** These alternatives would deeply impact the Execution Environment (Section 5) and Sharding (DSCs, Section 2.2), potentially making DSCs more akin to these ELs.
- **Comparison with Existing Implementations:**
  - **Ethereum L2 Rollups (Arbitrum, Optimism, zkSync):** The “Hybrid DAG + Execution Layer” is directly analogous to how L2s use Ethereum L1 for settlement and data availability. SNN would use its own DAG for this.
  - **Polkadot/Cosmos:** Parachains/Zones are somewhat similar to ELs, with the Relay Chain/Hub providing shared security and interoperability. SNN’s DAG would serve a similar role.
  - **SNN’s Approach:** The initial design (1.1-1.3) aims for high base-layer scalability. These alternatives are future considerations if even greater

specialization or capacity is needed, leveraging the DAG's strengths as a secure and decentralized ordering/settlement layer.

---

## 2. Scalability Architecture

This section details how SNN achieves scalability beyond the core DAG, through parallel processing inherent in the DAG, sharding via Dynamic Synaptic Clusters (DSCs), and efficient state management.

### 2.1. Parallel Processing

How the implementation of the DAG and transaction processing (as described in Section 1) achieves parallel processing.

- **Inherent DAG Parallelism:**
  - **Simultaneous TU Creation:** Multiple validators can create and issue new TUs concurrently without waiting for a global leader or block proposal slot. Each TU references existing tips, contributing to the DAG's parallel growth.
  - **Independent Transaction Validation:** Transactions affecting different account-synapses or state regions can be validated and included in TUs in parallel by different validators.
  - **Avalanche Consensus Parallelism:**
    - *Concurrent Queries:* Validators perform Avalanche queries for different TUs simultaneously.
    - *Localized Decisions:* A validator's decision to prefer/accept a TU is based on sampling a small subset of the network, not requiring global communication for each step.
- **Account-Synapse Model (from 1.1.1):**
  - **Concept:** Each user account effectively has its own "chain" of TUs (its account-synapse). Transactions that only affect a single account's state (e.g., internal transfers, simple state updates) can be processed and ordered along that account-synapse largely independently of others.
  - **Implementation:** TUs must reference the latest TU from the sender's account-synapse (enforcing per-account sequentiality). When a TU involves multiple accounts (e.g., a transfer), it references the latest TUs from all involved account-synapses, creating a join point in the DAG.
  - **Benefit:** This allows for a high degree of parallelism for transactions that don't have shared state contention.
- **Specific Algorithms & Data Structures Enabling Parallelism:**
  - **DAG Data Structure:** Allows multiple "heads" or "tips" to exist and be built upon simultaneously.

- **Mempool:** Can be structured to allow validators to pick transactions for different account-synapses or non-conflicting state regions in parallel.
- **Execution Environment (Section 5.1):** Must be designed to execute non-conflicting transactions (identified by DAG structure and account-synapses) in parallel post-consensus. This might involve techniques like Software Transactional Memory (STM) or an actor-based model.
- **Protocol Details & Message Flows:**
  - **TU Gossip:** TUs are gossiped independently.
  - **Consensus Queries:** Avalanche queries for different TUs proceed in parallel.
- **Component Breakdown:**
  - **Validator Client Software:** Must be multi-threaded to handle concurrent TU creation, validation, gossiping, and consensus participation.
  - **Execution Environment:** Key for realizing post-consensus execution parallelism.
- **Implementation Challenges & Trade-offs:**
  - **Challenge (State Contention):** While many transactions can be parallelized, those touching the same state (e.g., a popular smart contract, an exchange) still require sequential execution or sophisticated concurrency control at the execution layer. The DAG helps order them, but execution can still be a bottleneck.
  - **Challenge (Cross-Account-Synapse Synchronization):** Ensuring consistent views when TUs join multiple account-synapses. The DAG structure itself handles this ordering.
  - **Trade-off (Granularity of Parallelism vs. Overhead):** Very fine-grained parallelism (e.g., per single transaction TU) might increase overhead compared to batching some related transactions into micro-blocks/TUs.
- **Security Considerations:**
  - No new major security issues beyond those for the DAG itself, but parallelism can make debugging and tracing issues more complex.
- **Integration Points:**
  - **Core DAG Structure (1.1):** Provides the foundation.
  - **Transaction Processing (1.3):** Leverages this parallelism.
  - **Execution Environment (5.1):** Must be capable of exploiting the parallel ordering provided by the DAG.
  - **DSCs (2.2):** Take parallelism to a coarser, network-level sharding.
- **Research & Prototyping Needs:**
  - Benchmarking the degree of parallelism achievable in practice based on different transaction workloads.



- Developing advanced parallel execution engines for smart contracts on a DAG.
- **Comparison with Existing Implementations:**
  - **Aptos (Block-STM), Sui (Narwhal/Bullshark & object model):** These systems are explicitly designed for massive parallel execution. SNN's account-synapse model and DAG structure provide a similar foundation for identifying non-conflicting operations. SNN's Avalanche consensus differs from their consensus approaches.
  - **Traditional Blockchains:** Offer limited parallelism, mostly at the block validation stage (e.g., validating multiple transactions within a block if they are independent). SNN aims for much higher parallelism at the transaction ordering and consensus stages.
  - **SNN Innovation:** The combination of the hybrid DAG, account-synapses, and Avalanche consensus as a holistic system for achieving high degrees of parallel processing from transaction submission through to consensus.

## 2.2. Sharding/Subnets: “Dynamic Synaptic Clusters” (DSCs)

DSCs are SNN's approach to horizontal scaling, allowing the network to partition itself into manageable, interconnected sub-networks.

### 2.2.1. Concept (Dynamic Formation, Validator Roles)

- **Concept:** DSCs are effectively shards or subnets within the SNN ecosystem. Each DSC can be thought of as running its own instance of the Synaptic Ledger (or a compatible variation), processing transactions and reaching local consensus for state within that DSC.
- **Dynamic Formation:**
  - **Algorithm (Formation Trigger):** DSCs can be formed based on:
    1. **Network Load:** If the main SNN DAG (or an existing DSC) consistently exceeds certain load thresholds (e.g., transaction throughput, state growth rate, validator resource usage), governance can trigger the creation of a new DSC.
    2. **Application Demand:** A specific application or a group of applications requiring high throughput or specialized execution environments can request/sponsor the creation of a dedicated DSC via a governance proposal.
    3. **Geographic Proximity:** Potentially, for latency-sensitive applications, though this is harder to enforce permissionlessly.
  - **Algorithm (Validator Assignment):**
    - Validators can choose to join specific DSCs based on their interests, hardware capabilities, or delegation preferences.

- A minimum stake and/or a separate bonding process might be required to join a DSC.
  - There needs to be a mechanism to ensure sufficient validator distribution across DSCs for security. This could involve incentives or random sampling from the global validator pool for bootstrapping new DSCs.
  - A global “beacon” DAG (the main Synaptic Ledger or a dedicated one) would manage the registry of DSCs and their validator sets.
- **Validator Roles:**
  - **DSC Validators:** Responsible for processing transactions, participating in consensus (e.g., Avalanche) within their assigned DSC, and maintaining the state of that DSC.
  - **Global Validators (Optional/Main DAG):** Some validators might only participate in the main Synaptic Ledger, which handles cross-DSC communication, global state snapshots, or overall network security. Alternatively, all validators could be required to monitor the main DAG while also serving one or more DSCs.
  - **Relayers/Connectors:** Specialized roles (can be validators) responsible for facilitating cross-DSC communication.
- **Data Structures:**
  - **DSCRegistry (Main DAG Contract) :** Stores metadata for each DSC: ID, validator set, state root, communication endpoints, specific parameters (e.g., if it’s a specialized execution DSC).
- **Protocol Details & Message Flows:**
  - **DSC Creation:** Governance proposal -> Validator opt-in/assignment -> DSC initialization (genesis TU for the DSC) -> Registration in DSCRegistry .
- **Component Breakdown:**
  - **DSC Management Module (on Main DAG).**
  - **Validator DSC Assignment/Staking Logic.**
  - **Per-DSC Consensus Engine & P2P Network.**
- **Implementation Challenges & Trade-offs:**
  - **Challenge (Secure Validator Allocation):** Ensuring each DSC has enough honest validators to prevent takeover, especially for newly formed or less popular DSCs. Random sampling and incentive adjustments are key.
  - **Challenge (Dynamic Reconfiguration):** Handling validators joining/leaving DSCs smoothly without disrupting operations.
  - **Trade-off (Number of DSCs vs. Security):** More DSCs mean more parallelism but potentially dilutes security per DSC if the total validator pool isn’t large enough.

- **Trade-off (Homogeneity vs. Heterogeneity of DSCs):** Homogeneous DSCs (all run standard SNN logic) are simpler. Heterogeneous DSCs (specialized VMs, different parameters) offer flexibility but increase complexity.
- **Security Considerations:**
  - **Single-DSC Corruption:** If a DSC's validator set is compromised, its state can be corrupted. Cross-DSC communication protocols must be resilient to this (e.g., not blindly trusting messages from a potentially corrupt DSC without further validation or collateral).
  - **Validator Centralization Risk within a DSC.**
- **Parameterization & Configuration:**
  - Min/max validators per DSC.
  - Stake requirements for DSC participation.
  - Load thresholds for triggering new DSC creation.
  - Epoch length for DSC validator rotation.
- **Integration Points:**
  - **Main SNN DAG/Consensus (for DSC registry and cross-DSC coordination).**
  - **Staking & Governance Modules.**
  - **Cross-DSC Communication Protocol (2.2.3).**
- **Research & Prototyping Needs:**
  - Simulation of dynamic DSC formation and validator allocation algorithms.
  - Incentive models for encouraging balanced validator distribution.
- **Comparison with Existing Implementations:**
  - **Ethereum L2s/Rollups:** Each L2 is like a DSC, but SNN DSCs are envisioned as more tightly integrated first-class citizens.
  - **Polkadot Parachains, Cosmos Zones, Avalanche Subnets:** These are all forms of sharding/app-chains. SNN DSCs draw inspiration from these, particularly Avalanche Subnets for the dynamic and potentially permissionless nature.
  - **SNN Innovation:** Focus on “dynamic” formation based on load/demand and tight synergy with the base Synaptic Ledger for communication.

### 2.2.2. Synergy with DAG (Routing, Cross-Referencing)

How DSCs and the main Synaptic Ledger (or a global beacon DAG) interact.

- **Routing:**
  - **Transaction to DSC:** Users (or their wallets/RPC providers) need to determine which DSC a transaction belongs to.

- For transactions interacting with state primarily within one DSC, they are routed directly to validators of that DSC.
  - The `DSCRegistry` on the main DAG can be queried to find the target DSC for a given application or address range.
- **Cross-DSC Transaction Routing:** Handled by the cross-DSC communication protocol (2.2.3).
- **Cross-Referencing:**
  - **Concept:** DSCs periodically “anchor” or “commit” their state (or a summary of it, like a state root or a hash of recent TUs) to the main Synaptic Ledger. This provides a global point of reference and enhances the security of DSCs by making their history globally visible and verifiable.
  - **Data Structure (Anchor TU):** A special TU on the main Synaptic Ledger created by a DSC (e.g., by a committee of its validators). It contains: DSC ID, epoch number, state root of the DSC, list of significant cross-DSC messages sent/received.
  - **Algorithm (Anchoring):**
    1. DSC reaches local finality on a batch of TUs / an epoch.
    2. DSC validators compute the state root.
    3. A threshold of DSC validators sign this state root.
    4. This signed commitment is submitted as an Anchor TU to the main Synaptic Ledger.
    5. Main Synaptic Ledger validators verify the signatures and include the Anchor TU.
- **Benefits of Synergy:**
  - **Shared Security:** The main DAG can provide an additional layer of security/finality to DSC states.
  - **Interoperability:** Anchoring enables trust-minimized cross-DSC communication.
  - **Global State View (Partial):** The main DAG provides a high-level view of the entire SNN ecosystem’s state via these anchors.
- **Implementation Challenges & Trade-offs:**
  - **Challenge (Anchoring Frequency):** Too frequent anchoring increases load on the main DAG. Too infrequent increases latency for cross-DSC operations that rely on these anchors.
  - **Challenge (Size of Anchors):** Keeping anchor TUs small while conveying enough information.
  - **Trade-off (Tight vs. Loose Coupling):** Tighter coupling (more frequent/detailed anchors) improves consistency but reduces DSC autonomy and scalability.
- **Security Considerations:**

- **Invalid Anchors:** The main DAG must validate the signatures on Anchor TUs. If a DSC is corrupted and submits a false anchor, it should be slashable.
- **Parameterization & Configuration:**
  - Anchoring frequency per DSC.
  - Number of DSC validator signatures required for an anchor.
- **Integration Points:**
  - **Main Synaptic Ledger Consensus (processes Anchor TUs).**
  - **DSC Local Consensus.**
  - **Cross-DSC Communication Protocol.**
- **Research & Prototyping Needs:**
  - Optimal data structures for Anchor TUs.
  - Gas/fee model for submitting anchors to the main DAG.
- **Comparison with Existing Implementations:**
  - **Rollups (Optimistic/ZK):** Regularly submit state roots/proofs to L1. SNN's anchoring is similar.
  - **Cosmos Hub (Replicated Security):** Consumer chains can leverage the Hub's validator set. SNN's main DAG provides a similar security backstop via anchoring.
  - **SNN Innovation:** The DAG-to-DAG anchoring mechanism, where both main ledger and DSCs are DAG-based, potentially allowing for more flexible and efficient cross-referencing.

### 2.2.3. Cross-Cluster Communication (Direct DAG Referencing, Milestone Anchoring, Decentralized Relayers)

Mechanisms for DSCs to interact with each other. This is crucial for composability.

- **Direct DAG Referencing (for Tightly Coupled DSCs or within a “Super-Cluster”):**
  - **Concept:** If two DSCs are very closely related or trust each other to a higher degree (e.g., run by the same validator set or application suite), TUs in one DSC could directly reference TUs (or their hashes) in another DSC as parents. This is complex and implies validators in one DSC are at least partially aware of the other's DAG.
  - **Use Case:** Niche, high-performance needs. Not for general inter-DSC communication.
  - **Challenge:** Scalability, complexity of maintaining partial views of other DAGs.
- **Milestone Anchoring & Proof-Based Communication (Standard Method):**

- **Concept:** This is the primary method, relying on the anchoring mechanism (2.2.2). To send a message or transfer an asset from DSC-A to DSC-B:
  1. Transaction on DSC-A initiates the cross-DSC action (e.g., “lock asset X for transfer to user U on DSC-B”). This action is included in a TU and finalized within DSC-A.
  2. This event is included in DSC-A’s next anchor TU submitted to the main Synaptic Ledger. The anchor contains DSC-A’s state root which proves this event occurred.
  3. DSC-B validators (or relayers serving DSC-B) observe the main Synaptic Ledger for relevant anchor TUs from DSC-A.
  4. Once DSC-A’s anchor TU is finalized on the main DAG, a relayer submits a proof (Merkle proof of the event against DSC-A’s anchored state root, plus proof of the anchor’s finality on the main DAG) to a bridge/gateway contract on DSC-B.
  5. The gateway contract on DSC-B verifies this proof. If valid, it triggers the corresponding action on DSC-B (e.g., “mint wrapped asset X for user U”).
- **Specific Algorithms & Data Structures:**
  - `GatewayContract` (per DSC) : Smart contract on each DSC to receive and verify incoming cross-DSC messages/proofs and dispatch local actions.
  - `Merkle Proofs` / `SMT Proofs` : To prove inclusion of events in anchored state roots.
  - `CrossDSCtxReceipts` : Stored on the main DAG or verifiable from it, confirming a message was sent and anchored.
- **Protocol Details & Message Flows:** As described above.
- **Decentralized Relayers:**
  - **Concept:** Relayers are off-chain actors (can be validators, users, or dedicated services) that monitor events on one DSC (or the main DAG) and submit proofs/messages to another DSC.
  - **Incentivization:** Relayers are paid fees (by users initiating cross-DSC transactions or from protocol emissions) for their service.
  - **Security:** Relayers don’t typically hold user funds directly in this model. Their role is to ferry information. The security relies on the proofs they submit being verifiable by the destination DSC’s gateway contract. Malicious relayers submitting invalid proofs would have them rejected. Censorship is mitigated by allowing anyone to be a relayer.
- **Component Breakdown:**
  - `GatewayContract` on each DSC.
  - Relayer client software.

- Proof generation and verification libraries.
- Cross-DSC message queues/event logs within each DSC.
- **Implementation Challenges & Trade-offs:**
  - **Challenge (Latency):** Cross-DSC operations involve finality on DSC-A, finality of anchor on main DAG, and then processing on DSC-B. This can take time.
  - **Challenge (Proof Complexity & Gas Costs):** Verifying proofs on the destination DSC can be computationally intensive if not optimized (e.g., using ZK proofs for batched cross-DSC messages, similar to 3.1.3).
  - **Challenge (Atomicity):** Achieving atomic cross-DSC transactions (either both sides complete or neither) is very hard. Usually, a two-phase commit or saga pattern is used, with compensation logic if one side fails.
  - **Trade-off (Security vs. Speed):** Bypassing the main DAG for direct DSC-to-DSC communication (if validators overlap or via a shared communication bus) could be faster but might have weaker security assumptions than main DAG-mediated communication.
- **Security Considerations:**
  - **Relayer Collusion/Failure:** While proofs are verified, relayers could censor messages or fail. A robust, decentralized relayer network is needed.
  - **Replay Attacks across DSCs:** Gateway contracts must have replay protection for messages.
  - **State Inconsistency during Failures:** Robust error handling and compensation logic.
- **Parameterization & Configuration:**
  - Relayer fees.
  - Timeout periods for cross-DSC operations.
  - Proof verification gas limits on gateway contracts.
- **Integration Points:**
  - **Anchoring Mechanism (2.2.2).**
  - **Main Synaptic Ledger.**
  - **Smart Contract platform on each DSC (for GatewayContracts).**
  - **Asset Bridging (Section 3):** Similar principles for cross-chain vs. cross-DSC.
- **Research & Prototyping Needs:**
  - Efficient proof systems for cross-DSC messages (e.g., batched ZK proofs).
  - Mechanisms for near-atomic cross-DSC swaps or contract calls.
  - Robust relayer incentive models.

- **Comparison with Existing Implementations:**
  - **IBC (Cosmos):** The gold standard for trust-minimized cross-blockchain communication. SNN's milestone anchoring and relayer model is very similar in principle to IBC, using the main SNN DAG as the "Hub" for DSCs.
  - **Avalanche Cross-Subnet Communication:** Uses "Avalanche Warp Messaging" which allows subnets to send messages to each other, validated by the P-Chain validators. SNN's approach is analogous.
  - **Polkadot XCMP/HRMP:** Allows parachains to exchange messages. SNN's model is similar, with the main DAG acting like the Relay Chain.
  - **SNN Innovation:** Leveraging the DAG structure for both the main ledger and DSCs, potentially allowing for more nuanced or efficient cross-referencing and proof mechanisms compared to purely block-based systems.

## 2.3. State Management

Handling the state of the SNN, especially given the high throughput goals and DAG structure.

### 2.3.1. State Rent/Storage Fees

Mechanisms to ensure users pay for the ongoing cost of storing their data on the blockchain.

- **Concept:** Users pay not only for transaction execution but also for the persistent storage their transactions/contracts consume on the ledger. This can be an upfront fee or a recurring fee.
- **Specific Algorithms & Data Structures:**
  - **Algorithm (Fee Calculation):**
    - *Upfront (One-time):* When a contract is deployed or state is created/expanded (e.g., new entry in a map), a fee is calculated based on the size of the data ( `bytes_stored * fee_per_byte` ). This fee "pre-pays" for storage for a certain duration or indefinitely (if the model is simpler).
    - *Recurring (Rent):* Periodically (e.g., every epoch or year), accounts/contracts are charged a fee based on the amount of state they occupy. If the fee cannot be paid (e.g., account has no funds, contract has no mechanism to pay), the state can be made dormant or archived (see 2.3.2).
  - **Data Structure (Account/Contract State):** Each piece of state (e.g., contract storage slot, account balance entry) needs metadata: `size_in_bytes` , `last_rent_paid_epoch` , `owner_account_for_rent` .



- **Algorithm (Rent Collection):** A system process (could be part of epoch transition logic, or triggered by special transactions) iterates through active state, calculates rent, and deducts it.
- **Protocol Details & Message Flows:**
  - **State Creation:** User submits tx -> Fee calculated (execution + storage) -> Tx processed, state stored, fee deducted.
  - **Rent Payment:** Epoch transitions -> System iterates state -> Rent deducted from owner's balance or a designated contract fund. If payment fails, state marked for potential hibernation/archival.
- **Component Breakdown:**
  - **Fee Calculation Module (integrated with Execution Environment).**
  - **State Accounting Module (tracks storage per account/contract).**
  - **Rent Collection System Process/Contract.**
- **Implementation Challenges & Trade-offs:**
  - **Challenge (Complexity of Rent Collection):** Iterating all state can be resource-intensive. Solutions: probabilistic collection, per-access rent payment, or requiring users to "top-up" rent proactively.
  - **Challenge (User Experience):** Recurring rent can be confusing or burdensome for users if not managed well by wallets/dapps.
  - **Challenge (Defining "Owner" for Rent):** For complex shared contracts, who pays?
  - **Trade-off (Upfront vs. Recurring):** Upfront is simpler but might overcharge for short-lived state or undercharge for perpetual state if not priced dynamically. Recurring is fairer but more complex. SNN might start with upfront and explore recurring.
  - **Trade-off (Strictness of Eviction):** How aggressively to archive unpaid state.
- **Security Considerations:**
  - **Griefing by Filling Storage:** If storage fees are too low, attackers could fill state.
  - **Rent Collection Oracle/Process Security:** The mechanism for charging rent must be secure and accurate.
- **Parameterization & Configuration:**
  - Fee per byte of storage.
  - Rent rate and collection frequency.
  - Grace periods for unpaid rent.
  - Exemptions (e.g., core protocol contracts).
- **Integration Points:**
  - **Execution Environment (calculates storage used).**

- **Fee Market (4.2).**
- **State Pruning/Archival (2.3.2).**
- **Account Model.**
- **Research & Prototyping Needs:**
  - Economic modeling of different state rent schemes.
  - Efficient algorithms for rent collection or alternative “pay-per-access” models.
- **Comparison with Existing Implementations:**
  - **EOS, Arweave (Upfront/Endowment):** EOS had state RAM markets. Arweave has a one-time fee for permanent storage.
  - **NEAR Protocol (Recurring Rent):** Contracts pay rent for storage. If they can't, state can be evicted.
  - **Ethereum (High Upfront Gas for SSTORE):** Effectively an upfront fee, but no explicit recurring rent, leading to state bloat concerns. EIP-4444 (history expiry) and EIP-4844 (blob storage) address aspects of this.
  - **SNN Innovation:** Integrating a fair and sustainable state rent model tailored to the DAG structure, potentially with mechanisms for contracts to self-fund their rent or for users to sponsor storage.

### 2.3.2. Hierarchical State Storage & Pruning (Hot, Warm, Cold/Archival)

Managing the lifecycle of state to keep active validator nodes performant.

- **Concept:** Not all state needs to be immediately accessible by validators for processing new transactions. State can be categorized by access frequency and moved to different storage tiers.
  - **Hot Storage:** Actively used state, recent state. Kept in memory or very fast SSDs by validators. Required for immediate transaction processing.
  - **Warm Storage:** Less frequently accessed state, but still potentially needed. Stored on slower SSDs/HDDs by validators or full nodes. Can be loaded into hot storage on demand with some latency.
  - **Cold/Archival Storage:** Very old or rarely accessed state (e.g., state for which rent hasn't been paid, historical transaction data beyond a certain age). Not kept by default by validators. Stored by archival nodes, decentralized storage networks (e.g., Arweave, Filecoin), or users themselves. Retrieving this state is slow and may require explicit user action/fees.
- **Specific Algorithms & Data Structures:**
  - **Algorithm (State Tiering Logic):** Based on LRU (Least Recently Used), access frequency counters, age of state, or explicit contract declarations (e.g., a contract can mark some data as archivable).

- **Data Structure (State Database):** Must support efficient querying for hot state and mechanisms to offload/retrieve from warm/cold tiers. (e.g., RocksDB with different column families for hot/warm, pointers to archival locations).
- **Algorithm (Pruning):**
  - *Validators:* Prune state that has moved to cold/archival tiers from their local storage. They only keep a commitment (e.g., Merkle root) to the full historical state.
  - *Full Nodes:* May keep warm state for longer periods.
  - *Archival Nodes:* Store the full history.
- **Protocol Details & Message Flows:**
  - **Accessing Warm/Cold State:** If a transaction needs state not in a validator's hot storage:
    1. Validator attempts to fetch from its warm storage.
    2. If not found, it may query peer full nodes or a decentralized storage network. This will incur latency. The transaction might be deferred or rejected if state cannot be retrieved in a timely manner.
    3. Users might be required to provide "witnesses" or proofs for state they want to access if it's known to be cold.
- **Component Breakdown:**
  - **State Database Manager (with tiering support).**
  - **Archival Node Network/Interface.**
  - **State Retrieval Protocol (for warm/cold state).**
- **Implementation Challenges & Trade-offs:**
  - **Challenge (Defining Tiering Rules):** Optimal rules can be application-dependent.
  - **Challenge (Ensuring Availability of Cold State):** Reliance on archival nodes or decentralized storage requires incentives and robustness.
  - **Challenge (Latency of Accessing Non-Hot State):** Can impact UX for certain operations.
  - **Trade-off (Validator Storage Cost vs. Access Latency/Complexity):** Aggressive pruning reduces validator costs but increases complexity of accessing older state.
- **Security Considerations:**
  - **Availability of Archived State:** If archival nodes go offline or lose data, that state might become inaccessible. Using robust decentralized storage solutions is key.
  - **Integrity of Retrieved State:** Proofs (e.g., Merkle proofs against an anchored state root) are needed when fetching from untrusted archival sources.

- **Parameterization & Configuration:**
  - Thresholds for moving state between tiers (age, access count).
  - Number of archival nodes to replicate data to.
  - Fees for retrieving archival state.
- **Integration Points:**
  - **State Rent (2.3.1):** Unpaid rent can trigger archival.
  - **Execution Environment (needs to handle state access latency).**
  - **P2P Network (for finding/retrieving state).**
  - **Decentralized Storage Solutions (e.g., IPFS, Arweave integration).**
- **Research & Prototyping Needs:**
  - Performance analysis of different state tiering strategies.
  - Incentive mechanisms for archival node operators.
  - Efficient protocols for on-demand state retrieval with proofs.
- **Comparison with Existing Implementations:**
  - **Ethereum (History/State Pruning):** Ethereum clients support various pruning modes (e.g., Geth's light sync, full-but-pruned). EIP-4444 proposes making historical data beyond a year optional for nodes.
  - **NEAR (Hibernation):** Contracts whose storage rent isn't paid can have their state hibernated (archived).
  - **Celestia (Data Availability Sampling):** Focuses on providing DA for rollups, allowing L2s to not store all data themselves. SNN archival nodes serve a similar DA role for old SNN state.
  - **SNN Innovation:** A native, hierarchical storage model deeply integrated with the DAG and state rent, aiming for sustainable long-term state management from the outset.

### 2.3.3. State Proofs (e.g., Vector Commitments, Merkle Trees)

Mechanisms for compactly representing and proving the state of the SNN or parts of it. Essential for light clients, cross-DSC communication, and stateless validation.

- **Specific Algorithms & Data Structures:**
  - **Primary Choice: Sparse Merkle Tree (SMT) or Merkle Patricia Trie (MPT):**
    - **SMT (e.g., Jellyfish Merkle Tree, Aergo SMT):**
      - *Pros:* Efficient for sparse state (where many possible keys have no value), good for proving non-membership. Conceptually simpler for some DAG state models.
      - *Cons:* Proofs can sometimes be larger than MPTs for dense areas.
    - **MPT (Ethereum-style):**

- *Pros:* Efficient for Ethereum-like account models. Well-understood.
  - *Cons:* Can be complex. Proof of non-membership is less direct.
- **SNN Choice:** An SMT is likely a better fit for a DAG-based system, especially with account-synapses that might lead to a sparser global state view. The exact SMT variant would be chosen based on performance and proof size benchmarks. BLAKE3 would be the hash function.
- **Alternative/Complementary: Vector Commitments (VCs) (e.g., Kate-KZG Commitments, IPA-based):**
  - **Concept:** Commit to an ordered vector of data (e.g., all account balances in a DSC, or all storage slots of a contract) such that one can create small proofs for individual elements or ranges.
  - **Pros:** Constant-size proofs (for Kate), efficient proof aggregation. Useful for proving properties of the entire state or large contiguous chunks.
  - **Cons:** May require trusted setup (Kate), updates can be more complex than SMTs for sparse changes. Less ideal for proving single, arbitrary key-value pairs compared to SMTs.
  - **SNN Use:** Could be used for DSC state root commitments or for specific data structures within smart contracts that benefit from VCs.
- **Protocol Details & Message Flows:**
  - **State Root Calculation:** After each TU (or batch of TUs in an epoch) is finalized, the SNN node updates its local SMT and recalculates the root hash. This root hash represents the entire state of the ledger (or DSC).
  - **Proof Generation:** When a light client or another DSC needs to verify a piece of state (e.g., an account balance, inclusion of a transaction), a full node generates a Merkle proof (for SMT) or VC proof from the relevant state root.
  - **Proof Verification:** The recipient verifies the proof against a known, trusted state root (e.g., one anchored on the main DAG).
- **Cryptographic Primitives:**
  - **Hash Function:** BLAKE3 (for SMTs).
  - **Elliptic Curve Pairings (for Kate Commitments if used):** e.g., BLS12-381.
- **Component Breakdown:**
  - **SMT/MPT Library:** Core component for state representation.
  - **Proof Generation Service (part of validator/full node).**
  - **Proof Verification Logic (in light clients, gateway contracts).**
- **Implementation Challenges & Trade-offs:**
  - **Challenge (SMT Performance):** SMT updates and proof generation can be bottlenecks if not highly optimized, especially with very large states.

- **Challenge (Proof Size):** While SMT proofs are logarithmic, for very deep trees or complex states, they can still be non-trivial. Aggregation techniques (e.g., ZK-SNARKs over Merkle proofs) can help.
- **Trade-off (SMT vs. MPT vs. VC):** Choice depends on the exact state model, desired proof properties, and performance characteristics. SMTs are a good general default.
- **Security Considerations:**
  - **Soundness of Proof System:** The chosen Merkle tree or VC scheme must be cryptographically sound.
  - **Implementation Bugs:** Errors in the proof generation/verification logic can break security.
- **Parameterization & Configuration:**
  - SMT tree depth (if fixed, though usually dynamic).
  - Cache sizes for SMT nodes.
- **Integration Points:**
  - **All parts of SNN that require state verification:** Light clients, bridges (Section 3), cross-DSC communication (2.2.3), stateless validation (2.3.4), governance.
  - **State Database (stores the SMT/MPT nodes).**
- **Research & Prototyping Needs:**
  - Benchmarking different SMT implementations (Jellyfish, etc.) with SNN's expected workload.
  - Exploring ZK-SNARKs for compressing/batching state proofs (e.g., proving multiple SMT paths with one small ZK proof).
- **Comparison with Existing Implementations:**
  - **Ethereum (MPT):** Uses Merkle Patricia Tries extensively.
  - **Cosmos SDK (IAVL Trees):** Uses a balanced Merkle tree (IAVL) similar to SMTs.
  - **Polkadot (Trie-based):** Uses a Substrate trie variant.
  - **Aptos/Sui (Jellyfish Merkle Tree / Sparse Merkle Trees):** Modern chains often opt for SMTs.
  - **SNN Innovation:** Applying a highly optimized SMT (likely BLAKE3-based) tailored for the hybrid DAG structure and account-synapses, and potentially combining it with VCs for specific use cases or ZK-proof aggregation for ultra-efficient state verification.

#### 2.3.4. Stateless Validation (Partial)

Reducing the need for validators to hold the entire current state to validate new TUs/transactions.

- **Concept:** Validators can validate new TUs/transactions using only the TU itself and “witnesses” (state proofs for all state accessed by the transactions in the TU). The validator only needs to know the latest state root. This significantly lowers the barrier to entry for validators.
  - **Full Stateless Validation:** Validators hold no state beyond the root. All TUs come with complete witnesses.
  - **Partial Stateless Validation (More Practical Initial Goal for SNN):** Validators hold hot state (2.3.2). TUs for which all accessed state is hot can be validated directly. For TUs accessing warm/cold state, they must come with witnesses. This is a hybrid model.
- **Specific Algorithms & Data Structures:**
  - **Algorithm (Witness Generation):** The entity creating the TU (e.g., a user’s wallet connected to a full node, or a validator itself if it has the state) pre-computes all storage accesses the TU’s transactions will make and generates SMT proofs for that state against the current known state root. These proofs form the witness.
  - **Data Structure (TU with Witness):** The TU structure is extended to include the witness data (a collection of SMT proofs).
  - **Algorithm (Stateless Validation Logic):**
    1. Validator receives TU with witness and a claimed pre-state root.
    2. Validator verifies that the claimed pre-state root is consistent with its known latest finalized state root (or a recent one).
    3. For each transaction in the TU:
      - It uses the SMT proofs in the witness to reconstruct the required pre-existing state values.
      - It executes the transaction against this reconstructed state.
      - It calculates the new state values and updates to the SMT (locally, in memory).
    4. After processing all transactions, the validator computes the post-state root. This post-state root is then part of what Avalanche consensus acts upon.
- **Protocol Details & Message Flows:**
  - User/Proposer creates TU, generates/attaches witness.
  - Validator receives TU, verifies witness against its state root, executes, derives new state root.
- **Component Breakdown:**
  - **Witness Generation Module (in clients/full nodes).**
  - **Stateless Execution Logic (in validators, integrated with Execution Environment).**
  - **SMT Proof Verification Library.**

- **Implementation Challenges & Trade-offs:**
  - **Challenge (Witness Size):** Witnesses can be large, increasing network bandwidth and TU size. Witness compression or ZK-ification is advanced research.
  - **Challenge (Predicting State Access):** Accurately predicting all state an arbitrary smart contract call will access can be hard (halting problem).  
Solutions:
    - Require transactions to declare accessed state keys (like Solana).
    - Execute once to find accesses, then generate witness (adds latency).
    - Restrict complex dynamic state access patterns.
  - **Challenge (State Contention with Witnesses):** If multiple TUs with witnesses try to modify the same state, their witnesses (based on an older state root) might become invalid. Requires careful ordering and witness regeneration or a model where witnesses prove against the state *after* parent TUs are applied.
  - **Trade-off (Statelessness vs. Witness Overhead/Complexity):** Full statelessness has high overhead. Partial statelessness is a good compromise.
- **Security Considerations:**
  - **Invalid Witnesses:** Validators must rigorously verify witnesses. A TU with an invalid witness (that doesn't match the claimed pre-state root, or is incomplete) must be rejected.
  - **Exploiting State Access Prediction:** If prediction is imperfect, it could be a vulnerability.
- **Parameterization & Configuration:**
  - Max witness size.
  - Rules for when witnesses are required (e.g., based on state tier).
- **Integration Points:**
  - **State Proofs (2.3.3):** Witnesses are made of these.
  - **Execution Environment (must support execution with witnesses).**
  - **Transaction Format (needs to accommodate witnesses).**
  - **Hierarchical Storage (determines when witnesses are needed).**
- **Research & Prototyping Needs:**
  - Efficient witness generation algorithms.
  - Techniques for witness compression or using ZK-proofs as witnesses.
  - Strategies for handling state access prediction for complex smart contracts.
  - Impact of witness propagation on network latency.
- **Comparison with Existing Implementations:**



- **Ethereum (Stateless Ethereum Research):** A long-term research goal for Ethereum. Concepts like “block witnesses” are similar. SNN can learn from this research.
  - **Mina Protocol (Succinct Blockchain):** Uses ZK-SNARKs to keep the chain constant size, effectively a form of full stateless validation where blocks themselves are proofs of transition. SNN’s approach is more about validator state, not necessarily a constant-size chain.
  - **Aptos/Sui:** While not fully “stateless” in the Ethereum sense, their designs aim to reduce validator state burden through efficient state models and proof systems.
  - **SNN Innovation:** Aiming for partial stateless validation from early on, integrated with the DAG structure and hierarchical storage, to lower validator requirements and enhance decentralization. The “account-synapse” model might simplify witness generation for many common transaction types.
- 

### 3. Interoperability with Ethereum & Bitcoin

SNN is designed for seamless and secure interaction with established networks like Ethereum and Bitcoin, facilitating asset movement and data exchange. This section details the technical approaches for bridging assets, passing data/messages, adhering to standards, and managing the inherent security trade-offs.

#### 3.1. Asset Bridging

SNN will employ a multi-layered approach to asset bridging, offering users a spectrum of options balancing security, trust assumptions, and speed. The primary goal is to allow SNN users to interact with assets from Ethereum and Bitcoin, and for SNN-native assets (like NEURON or SNN-based tokens) to be represented on these external chains.

##### 3.1.1. Light-Client Based Bridges (SNN <-> Ethereum)

This approach offers high security by enabling on-chain verification of the counterparty chain’s state through light clients.

- **Concept:**
  - **SNN verifying Ethereum:** SNN validators (or a designated, staked set of them) collectively run an Ethereum light client. They process Ethereum block headers, verify PoW/PoS attestations, and track relevant contract states (e.g., lock contracts) on Ethereum.

- **Ethereum verifying SNN:** An Ethereum smart contract acts as an SNN light client. It verifies SNN TUs that serve as “milestones” or “anchors” (containing SNN state roots from Section 2.2.2 and 2.3.3), their SMT proofs, and the collective signatures of SNN validators, effectively tracking the SNN chain’s state.
- **Specific Algorithms & Data Structures:**
  - **SNN-side Ethereum Light Client Service:**
    - **Data Structures:** `EthereumBlockHeaderStore` (validated headers), `EthereumLogMonitor` (tracks lock/unlock events from specific Ethereum contracts via Merkle Patricia Trie proofs).
    - **Algorithms:** Ethereum header validation (Ethash PoW or Casper FFG/LMD GHOST PoS), MPT proof verification.
  - **Ethereum-side SNN Light Client Contract:**
    - **Data Structures:** `SNNAnchorTUSore` (validated SNN anchor TU hashes and state roots), `SNNValidatorSetManager` (manages SNN validator public keys/stakes for signature verification, updated via SNN governance actions proven to this contract).
    - **Algorithms:** SNN anchor TU validation (verifying a supermajority of SNN Ed25519 signatures on the anchor TU), SMT proof verification for SNN state (e.g., proving an asset lock on SNN).
- **Protocol Details & Message Flows:**
  - **ETH/ERC20 from Ethereum to SNN (Lock & Mint):**
    1. User locks ETH/ERC20 in a `LockingContract` on Ethereum, emitting `LockedAsset(user_snn_address, token_address, amount)`.
    2. SNN validators’ Ethereum light client service observes and validates this event.
    3. A quorum of SNN validators submits this event proof (Ethereum block header, MPT proof of log) to a `BridgeOracleContract` on SNN.
    4. `BridgeOracleContract` (SNN) verifies proof, instructs `WrappedAssetContract` (SNN) to mint wrapped ETH/ERC20 (e.g., sETH, sDAI) to `user_snn_address`.
  - **Wrapped SNN Assets from SNN to Ethereum (Burn & Unlock):**
    1. User calls `burnWrappedAsset(ethereum_target_address, amount)` on `WrappedAssetContract` (SNN). This burns the asset and creates a `BurnEvent` TU (or includes it in an anchor).
    2. A relayer observes this, fetches the SNN anchor TU proof (validator signatures, SMT proof of the burn event against the anchor’s state root).
    3. Relayer submits proof to `SNNLightClientContract` on Ethereum.
    4. Ethereum contract verifies SNN anchor and burn proof, instructs `LockingContract` (ETH) to unlock original ETH/ERC20.

- **Cryptographic Primitives:**
  - **SNN:** Ed25519 (validator signatures), BLAKE3 (SMTs, TU hashes).
  - **Ethereum:** ECDSA (secp256k1), Keccak256 (MPT, block hashes), Ethash/PoS primitives.
- **Component Breakdown:**
  - **SNN Chain:** `EthereumLightClientService` (in SNN validator clients), `BridgeOracleContract` (SNN), `WrappedAssetContract` (SNN).
  - **Ethereum Chain:** `SNNLightClientContract` (ETH), `LockingContract` (ETH).
  - **Relayers:** Off-chain, permissionless, incentivized.
- **Implementation Challenges & Trade-offs:**
  - **Gas Costs on Ethereum:** Verifying SNN anchor TUs (many signatures) and SMT proofs on Ethereum is expensive. ZK-Rollups for aggregation (3.1.3) are crucial.
  - **Complexity:** Maintaining light clients for evolving chains.
  - **Latency:** Finality on both chains + relaying.
  - **Validator Set Updates:** Securely updating SNN validator set info on the Ethereum contract.
  - **Trade-off:** High security vs. complexity/cost.
- **Security Considerations (Implementation Specific):**
  - Light client bugs. Ethereum reorgs handling by SNN. SNN consensus failures affecting Ethereum contract. Relayer censorship (mitigated by permissionless relaying).
- **Parameterization & Configuration:**
  - Confirmation depths (ETH & SNN). Relayer fees. Gas limits.
- **Integration Points:** SNN Consensus (1.1), SMTs (2.3.3), SNN Validators, Ethereum Network.
- **Research & Prototyping Needs:**
  - Gas-efficient SNN anchor/SMT verification in Solidity (potentially via ZK pre-compiles).
  - Secure SNN validator set update mechanism on Ethereum.
- **Comparison with Existing Implementations:** Rainbow Bridge (NEAR), Gravity Bridge (Cosmos), IBC. SNN aims for similar trust-minimization, leveraging its DAG structure for efficient anchor generation.

### 3.1.2. MPC-Based Bridges (Bitcoin & Broader Compatibility)

For chains like Bitcoin lacking rich smart contract capabilities, or as a faster, more versatile (but more trust-reliant) option.

- **Concept:** A committee of SNN validators (or a dedicated, heavily staked MPC group within the SNN ecosystem) uses Multi-Party Computation (MPC) and Threshold Signature Schemes (TSS) to collectively manage private keys for addresses on external chains (e.g., a Bitcoin multisig or a single address where the key is sharded). No single entity holds the complete key.
- **Specific Algorithms & Data Structures:**
  - **MPC/TSS:**
    - **Algorithm:** Threshold ECDSA (e.g., GG18/GG20) for Bitcoin/Ethereum. Threshold Schnorr (e.g., FROST) if applicable.
    - **Protocols:** Distributed Key Generation (DKG) for committee key shares. Interactive signing protocol (  $t$ -of- $n$  threshold).
  - **Data Structures (SNN Contracts):**
    - `MPCCommitteeRegistry` : Manages MPC members, public key shares,  $t$ ,  $n$  values.
    - `ExternalChainAddressStore` : Publicly lists MPC-controlled addresses on BTC, ETH, etc.
    - `PendingOutgoingTxQueue` : For bridging assets out of SNN via MPC.
- **Protocol Details & Message Flows:**
  - **Asset from External Chain (e.g., BTC) to SNN (Deposit & Mint):**
    1. User sends BTC to an MPC-controlled Bitcoin address.
    2. MPC committee members (acting as watchtowers) observe this deposit.
    3. A quorum of MPC members attests to this deposit on SNN (to `BridgeOracleContract` ).
    4. `BridgeOracleContract` verifies attestations, instructs `WrappedAssetContract` to mint sBTC.
  - **Wrapped SNN Asset (e.g., sBTC) to External Chain (Burn & Sign):**
    1. User burns sBTC on SNN, specifying destination Bitcoin address. Request queued in `PendingOutgoingTxQueue` .
    2. MPC committee members pick up request, construct Bitcoin transaction.
    3. Engage in  $t$ -of- $n$  MPC signing protocol for the Bitcoin transaction.
    4. One or more members broadcast the signed Bitcoin transaction.
- **Cryptographic Primitives:**
  - Threshold ECDSA/Schnorr. Standard ECDSA/Schnorr for external chains. Ed25519 for SNN attestations.
- **Component Breakdown:**

- **SNN Chain:** MPCCommitteeManagerContract , BridgeOracleContract , WrappedAssetContract , PendingOutgoingTxQueue .
- **MPC Committee Nodes:** SNN nodes with MPC software, external chain clients.
- **Implementation Challenges & Trade-offs:**
  - **MPC Protocol Complexity & Security:** DKG/TSS are hard to implement correctly.
  - **Liveness & Key Share Security:**  $t$  members must be online and secure their shares.
  - **Scalability of MPC Signing:** Can be communication-heavy.
  - **Trade-off:** Broader compatibility (Bitcoin) & speed vs. trust in MPC committee's honesty/security (mitigated by SNN stake slashing).
- **Security Considerations (Implementation Specific):**
  - MPC software bugs. Collusion of  $t$  MPC members (primary risk). DoS against committee.
- **Parameterization & Configuration:**
  - $t$  ,  $n$  for MPC. MPC stake requirements. Confirmation depths. Key share refreshment frequency.
- **Integration Points:** SNN Staking/Slashing, SNN smart contracts, External chain networks.
- **Research & Prototyping Needs:**
  - Secure implementation and benchmarking of chosen TSS. Robust DKG/refresh protocols.
- **Comparison with Existing Implementations:** WBTC (centralized custodian, MPC is far more decentralized), RenBridge (used MPC, SNN learns from it), Thorchain (uses TSS). SNN's MPC bridge is a protocol-native, staked solution.

### 3.1.3. ZK-Rollups for Bridge State Aggregation

To reduce costs (especially gas on Ethereum for the light-client bridge) and improve efficiency of verifying SNN state (or batches of bridge operations) on other chains.

- **Concept:** Instead of SNN light client contract on Ethereum verifying many SNN anchor TU signatures and SMT proofs individually, a ZK-SNARK/STARK is used. A prover generates a single, small ZK proof attesting to the validity of a batch of SNN state transitions relevant to the bridge (e.g., many BurnEvent TUs for SNN->ETH transfers). This single proof is then verified by a simpler, cheaper verifier contract on Ethereum.
- **Specific Algorithms & Data Structures:**

- **ZK-SNARK/STARK System:**
  - **Choice:** Plonk/Groth16 (SNARKs) or STARKs. Lean towards universal/updatable trusted setups (e.g., Plonk with KZG) or transparent setups (STARKs).
  - **Circuit Design:** ZK circuit to:
    1. Verify SNN validator signatures on a sequence of anchor TUs.
    2. Verify SMT proofs for bridge-related events (e.g., burn events) against state roots in those anchor TUs.
    3. Aggregate results (e.g., total amounts for asset unlocks on Ethereum).
- **Data Structures (SNN):** `BridgeEventBatcher` , `ZKProofGeneratorService` (off-chain, incentivized provers).
- **Data Structures (Ethereum):** `SNN_ZKVerifierContract` (verifies ZK proofs, updates `LockingContract` ).
- **Protocol Details & Message Flows (Example: SNN -> Ethereum Batched Unlocks via ZK Proof):**
  1. `BurnEvent` TUs collected by `BridgeEventBatcher` on SNN.
  2. `ZKProofGeneratorService` takes a batch, generates  $\pi_{zk}$  attesting to their validity and aggregated unlock data.
  3.  $\pi_{zk}$  + public inputs submitted to `SNN_ZKVerifierContract` on Ethereum.
  4. `SNN_ZKVerifierContract` verifies  $\pi_{zk}$  . If valid, instructs `LockingContract` to perform batched unlocks.
- **Cryptographic Primitives:** Chosen ZK-SNARK/STARK construction (e.g., Plonk with KZG on BLS12-381).
- **Component Breakdown:**
  - **SNN Chain:** `BridgeEventBatcher` .
  - **Off-Chain (SNN Ecosystem):** `ZKProofGeneratorService` (decentralized provers).
  - **Ethereum Chain:** `SNN_ZKVerifierContract` , `LockingContract` .
- **Implementation Challenges & Trade-offs:**
  - **ZK Circuit Complexity & Auditing:** Highly specialized, error-prone.
  - **Proving Time & Cost:** Computationally intensive.
  - **Trusted Setup (some SNARKs).**
  - **Trade-off:** Massively reduced Ethereum gas costs & better scalability for bridge operations vs. high R&D, proving costs, potential prover centralization if not incentivized well.
- **Security Considerations (Implementation Specific):**
  - Soundness bugs in ZK circuit/prover (critical). Prover censorship.

- **Parameterization & Configuration:** Batch size, proof frequency, prover incentives.
- **Integration Points:** Augments light-client bridge (3.1.1), SNN SMTs (2.3.3).
- **Research & Prototyping Needs:**
  - Selection/benchmarking of ZK schemes. ZK circuit development for SNN bridge logic. Decentralized prover network design.
- **Comparison with Existing Implementations:** L2 Rollups on Ethereum (zkSync, StarkNet). SNN uses ZK for bridge efficiency, not full L2 scaling on top of Ethereum. Succinct/ZK bridges (e.g., some Mina concepts, Polygon Hermez).

### 3.1.4. Security Mechanisms for Bridged Assets

General security measures applicable across all SNN bridge types.

- **Asset Representation:** Wrapped tokens on SNN (sETH, sBTC, etc.). Original locked on native chain.
- **Total Value Locked (TVL) Monitoring & Global Rate Limits:**
  - On-chain monitoring of TVL in bridge contracts.
  - Governance-controlled global and per-asset rate limits on bridging velocity (amount per hour/day) to mitigate rapid drains.
  - Emergency thresholds triggering cooldowns or requiring Security Council approval for large, anomalous withdrawals.
- **Emergency Pause/Shutdown (Circuit Breaker):**
  - Triggered by SNN Governance (NEURON vote) and/or a time-locked SNN Security Council (elected by NEURON holders, acts faster in emergencies).
  - Can halt all bridge operations. Focus on enabling safe withdrawal of original assets based on last known valid state.
- **Validator Slashing for Bridge Misconduct:**
  - SNN validators in light-client verification groups or MPC committees face severe slashing of staked NEURON for proven malicious actions (e.g., attesting false external events, MPC collusion). Proofs submitted to SNN's slashing contract.
- **Decentralized Relayer/Prover Networks:**
  - For light-client and ZK-rollup bridges, ensure relayers/provers are decentralized, permissionless (with bonds/incentives), and fault-tolerant. Users can self-relay/prove.
- **Insurance Funds (Optional):**
  - Portion of bridge fees to an insurance fund (governed by NEURON holders) to potentially compensate users for losses from *unforeseen*

bridge bugs (prevention is primary).

- **Rigorous Audits & Bug Bounties:**
  - Mandatory, regular security audits for all bridge components (SNN contracts, external chain contracts, MPC software, ZK circuits) by multiple reputable firms.
  - Generous, ongoing bug bounty program.
- **Transparency & User Warnings:** Clear communication of security models, trust assumptions, TVL, rate limits for each bridge.
- **Implementation Challenges:** Designing non-gameable rate limits/emergency mechanisms. Secure Security Council governance. Correct attribution for slashing.
- **Integration Points:** SNN Governance (6.2), Staking/Slashing (6.1).

### 3.2. Data/Message Passing

Enabling SNN to exchange arbitrary data or messages with other chains, not just assets.

- **Decentralized Oracle Network (DON) - Integration or Fostering:**
  - **Concept:** SNN itself doesn't aim to be a full-blown oracle network initially. However, it will:
    1. **Integrate with Existing DONs:** Provide easy ways for SNN smart contracts to consume data from established DONs like Chainlink, Band Protocol, etc. This involves SNN contracts that can verify oracle reports originating from these DONs on other chains (e.g., Ethereum). These reports would typically be relayed to SNN and verified against the DON's on-chain commitments on their native chain.
    2. **Foster SNN-Native Oracles:** The SNN platform (especially with DSCs) could become an attractive environment for new oracle solutions to be built. SNN governance might incentivize this.
  - **Specific Algorithms & Data Structures:**
    - `OracleAggregatorContract (SNN)` : A contract that receives data reports from multiple oracle nodes (either SNN-native or relayed from external DONs), verifies their signatures/stakes, and computes an aggregate result (e.g., median).
    - Relayers for bringing external DON reports onto SNN.
  - **Security:** Relies on the security of the chosen DON. For SNN-native oracles, stake-based security and reputation would be key.
- **Cross-Chain Invocation via Bridges (Generic Messaging):**
  - **Concept:** Extend the asset bridging mechanisms (especially light-client or ZK-rollup based) to support passing generic messages.



- An SNN contract calls a `GatewayContract` on SNN, specifying a target contract and function call data for an Ethereum contract.
- This message is anchored on the main SNN DAG (if originating from a DSC) and then relayed to the `SNNLightClientContract` (or `SNN_ZKVerifierContract`) on Ethereum.
- The Ethereum contract, after verifying the SNN message's authenticity, can then perform a `CALL` to the target Ethereum contract with the provided data.
- Return values can be passed back via a similar reverse mechanism.
- **Challenges:** Asynchronous nature, gas costs on target chain, ensuring atomicity or clear handling of failures is complex.
- **State Proof Verification (as a form of Data Passing):**
  - **Concept:** SNN contracts can be made aware of specific state on Ethereum (or other chains with light clients on SNN) by verifying Merkle proofs of that state.
    - Example: An SNN contract might need to verify an account's NFT ownership on Ethereum before unlocking features on SNN.
    - Relayers fetch the required state proof from Ethereum and submit it to the SNN contract, which uses the SNN-side Ethereum light client service (3.1.1) to verify the block header and then the MPT proof.
  - **Security:** Relies on the security of SNN's Ethereum light client.
- **Component Breakdown:**
  - `GatewayContract` (SNN, for outgoing generic messages).
  - `SNNLightClientContract` (Ethereum, extended for incoming generic messages).
  - Oracle consumer contracts/libraries for SNN.
- **Implementation Challenges:**
  - Standardizing message formats for cross-chain calls.
  - Managing gas payments for invocations on the target chain.
  - Handling asynchronicity and lack of atomicity gracefully.
- **Comparison with Existing Implementations:** LayerZero, Axelar, Wormhole (generic messaging bridges). IBC (generic messaging in Cosmos). SNN aims for a similar capability, prioritizing security via light clients / ZK proofs where possible.

### 3.3. Standardization

Adopting and adapting existing cross-chain standards to maximize compatibility.

- **Focus on Specific Technical Standards:**

- **IBC (Inter-Blockchain Communication Protocol):**
  - **Adaptation:** While SNN is not a Cosmos SDK chain, the principles of IBC (light clients, relayers, connection/channel handshakes, packet semantics) are highly relevant. SNN will aim to implement an IBC-compatible interface, allowing it to connect to the broader IBC ecosystem. This involves:
    - Implementing SNN's light client logic in a way that can be consumed by IBC relayers and other IBC-enabled chains.
    - Building modules on SNN that can understand and process IBC packets.
    - This would likely involve an "IBC Core" module on SNN and specific application modules (like token transfer) that use it.
- **CCIP (Chainlink Cross-Chain Interoperability Protocol):**
  - **Integration:** As CCIP matures and sees wider adoption, SNN will look to integrate as a supported network. This would involve SNN validators potentially participating in CCIP's DONs (if applicable) or SNN contracts being able to send/receive messages via CCIP routers.
- **ERC Standards (Ethereum):** For wrapped assets on SNN representing Ethereum tokens, adhere to ERC-20 (fungible), ERC-721/1155 (NFTs) interfaces to ensure compatibility with SNN wallets and dApps. Similarly, SNN-native assets bridged to Ethereum should be represented as ERC-compatible tokens.
- **CAIPs (Chain Agnostic Improvement Proposals):** Adhere to relevant CAIPs for chain IDs (CAIP-2), asset IDs (CAIP-19), account IDs (CAIP-10) to simplify integration with multi-chain wallets and services.
- **Implementation Challenges:**
  - Significant engineering effort to implement and maintain compatibility with complex protocols like IBC.
  - Keeping up with evolving standards.
- **Benefits:** Greatly enhances SNN's reach and utility by tapping into existing networks and liquidity.

### 3.4. Security and Decentralization Trade-offs

Analyzing the trade-offs for each bridge type and messaging approach.

- **Light-Client Bridges (e.g., SNN <-> Ethereum):**
  - **Security:** Highest (trust-minimized), relies only on the honesty of the two chains' consensus and correctness of light client implementations.
  - **Decentralization:** High, as it relies on the full validator sets of both chains.
  - **Trade-offs:** Highest complexity, highest on-chain gas costs (especially on Ethereum side), potentially higher latency due to finality

requirements.

- **MPC-Based Bridges (e.g., for Bitcoin):**
    - **Security:** Relies on the honesty of  $t$ -of- $n$  MPC committee members and the security of the MPC protocols. Less trust-minimized than light clients.  $(n-t+1)$  collusion can break it.
    - **Decentralization:** Depends on the size ( $n$ ) and distribution of the MPC committee. Can be highly decentralized if  $n$  is large and diverse, but still a smaller set than all SNN validators.
    - **Trade-offs:** Lower on-chain gas costs (no complex verification on external chain), broader compatibility (works with non-smart contract chains), potentially faster. The main trade-off is introducing this specific honest-majority assumption for the MPC committee. (Mitigated by SNN stake slashing for MPC members).
  - **ZK-Rollups for Bridge State Aggregation:**
    - **Security:** Relies on the soundness of the ZK cryptography and correctness of the circuit. If sound, offers security comparable to light clients but for batched operations. Potential risk from bugs in complex ZK circuits or prover issues.
    - **Decentralization:** Prover network decentralization is a concern. If provers are centralized, they can censor, though not steal funds if ZK math is sound.
    - **Trade-offs:** Significantly reduces gas costs for light-client style verification. High R&D complexity. Proving can be slow/costly.
  - **Generic Messaging (e.g., LayerZero-style or via SNN's bridges):**
    - **Security:** Varies greatly by provider. Some rely on multi-sig oracles, others on light clients (like IBC). SNN's native generic messaging via its light-client/ZK bridges would inherit their security properties. Using external providers introduces their trust assumptions.
    - **Decentralization:** Varies by provider.
    - **Trade-offs:** Convenience and broad reach vs. potentially weaker trust assumptions than SNN's native bridges.
  - **SNN's Approach:** Offer a spectrum. Prioritize light-client/ZK bridges for core asset flows where feasible (especially with Ethereum). Use MPC for Bitcoin and as a pragmatic alternative. Clearly communicate the trade-offs of each option to users and developers. Governance can play a role in sanctioning/auditing bridge components.
- 

## 4. User Experience (UX) Revolution

SNN aims to significantly improve blockchain UX through native account abstraction, predictable low fees, human-readable addressing, and simplified onboarding.

## 4.1. Account Abstraction Native

Implementing smart contract accounts as the default, offering flexibility like social recovery, multi-sig, session keys, and paying fees in various tokens.

- **Concept:** Every user account on SNN is a smart contract (an “Account Contract” or AC). Externally Owned Accounts (EOAs) as seen in Ethereum do not exist in the same way. Users interact with their ACs, which then interact with other contracts or ACs. This allows account logic to be programmable.
- **Specific Algorithms & Data Structures:**
  - **Data Structure (Account Contract - AC):**
    - Standardized interface (e.g., `I SNN_Account`):
      - `validateTransaction(transaction)` : Validates a transaction intended for this AC (checks signature, nonce, permissions). This is the core of AA.
      - `executeTransaction(transaction)` : Executes the validated transaction.
      - `getNonce()`
      - Functions for managing signers, recovery mechanisms, session keys, per-app allowances, preferred fee token.
    - Storage: Stores signers (can be multiple, weighted), social recovery partners, session key configurations, per-app allowances, preferred fee token.
  - **Algorithm (Transaction Validation Flow):**
    1. User (via wallet) crafts a “UserOperation” (similar to EIP-4337) specifying target, data, gas limits, signature, etc.
    2. This UserOperation is sent to SNN validators (acting as “Bundlers” in EIP-4337 terms, or it’s just a standard SNN transaction targeting the AC).
    3. The SNN protocol ensures that before executing the UserOperation’s main action, the target AC’s `validateTransaction()` method is called with the UserOperation.
    4. `validateTransaction()` within the AC implements custom logic:
      - Verifies the signature against its stored signers (can be EdDSA, ECDSA via precompile, or even quantum-resistant sigs if the AC supports it).
      - Checks nonce.

- Checks if the operation is permitted by session keys (if applicable).
- Can implement logic for fee payment in tokens other than NEURON (see 4.1.x).

5. If `validateTransaction()` succeeds, the main operation is executed.

- **Data Structure (Global Entry Point / Singleton):** SNN might have a global “EntryPoint” contract (similar to EIP-4337) that all `UserOperations` are sent to. This `EntryPoint` then calls `validateTransaction` on the target AC. Alternatively, the SNN protocol itself can enforce this call to the AC directly. A direct protocol-level enforcement is more “native.”
- **Implementation Details for Features:**
  - **Smart Contract Accounts:** As described. Each AC has its own address.
  - **Social Recovery:** AC stores a list of trusted “guardian” AC addresses. A recovery process (e.g., m-of-n guardians agree) can change the primary signer(s) of the AC.
  - **Multi-sig:** AC’s `validateTransaction` can require multiple signatures from different stored signers.
  - **Session Keys:** AC can grant temporary, restricted permissions (e.g., for a specific game contract, limited transaction value/count) to a short-lived key. `validateTransaction` checks these session key rules.
  - **Fee Payment in Various Tokens:**
    - AC’s `validateTransaction` can authorize a “Paymaster” contract (see 4.2.3) to pay fees in NEURON on its behalf, in exchange for the AC transferring other tokens (e.g., sUSDC) to the Paymaster.
    - Alternatively, the SNN protocol could natively support fee payment in a whitelist of tokens, with on-chain DEX integration for conversion to NEURON for validator payment.
  - **Batch Transactions:** AC’s `executeTransaction` can be designed to take an array of operations and execute them sequentially. `validateTransaction` would validate the batch.
- **Cryptographic Primitives:**
  - Flexible: ACs can define their own signature schemes. EdDSA (for NEURON-based operations like staking, if ACs can be validators) and ECDSA (for Ethereum compatibility) should be supported via efficient precompiles.
  - Hash functions (BLAKE3).
- **Component Breakdown:**
  - **Account Contract Standard Library/Templates:** Pre-audited AC templates for common use cases (simple signer, multi-sig, social recovery).

- **EntryPoint Contract (Optional, if not protocol-enforced).**
- **Precompiles for Signature Verification (EdDSA, ECDSA, etc.).**
- **Wallet SDKs supporting UserOperations and AC management.**
- **Implementation Challenges & Trade-offs:**
  - **Challenge (Gas Overhead):** AC logic (especially `validateTransaction`) adds gas cost to every transaction compared to simple EOA signature checks. Optimization is key.
  - **Challenge (Complexity for Users/Wallets):** While powerful, managing AC features requires good wallet support.
  - **Challenge (Standardization):** Ensuring a secure and flexible AC interface standard.
  - **Trade-off (Native AA vs. EIP-4337 style):** Native AA (protocol enforces AC validation) is more seamless but harder to implement in the base protocol. EIP-4337 style (EntryPoint contract) is an overlay, easier to add but might have centralization points (Bundlers, though SNN validators can be bundlers). SNN aims for NATIVE AA.
- **Security Considerations (Implementation Specific):**
  - **Bugs in AC Implementation:** A bug in a user's AC could lead to loss of funds. Audited templates are crucial.
  - **Replay Attacks across different ACs/Chains:** Nonce management must be robust per AC.
  - **Denial of Service by Complex `validateTransaction`:** Gas limits on `validateTransaction` are essential.
- **Parameterization & Configuration:**
  - Gas costs for AC operations and precompiles.
  - Default AC templates.
- **Integration Points:**
  - **Transaction Processing (1.3):** Needs to understand UserOperations and trigger AC validation.
  - **Execution Environment (5.1):** Executes AC logic.
  - **Fee Market (4.2):** Interacts with fee payment logic in ACs.
  - **Wallet Ecosystem.**
- **Research & Prototyping Needs:**
  - Gas optimization for common AC operations.
  - Secure and user-friendly designs for social recovery and session keys.
  - Standard AC interface definition.
- **Comparison with Existing Implementations:**
  - **Ethereum (EIP-4337):** SNN provides native AA, avoiding EIP-4337's need for separate Bundler/Paymaster infrastructure (though Paymasters can

still exist as contracts). SNN validators handle bundling.

- **StarkNet, zkSync Era (Native AA):** These L2s have native AA. SNN learns from their designs, implementing it on a L1 DAG.
- **Argent, Gnosis Safe (Smart Wallets on Ethereum):** These are smart contract wallets built on top of Ethereum's EOA model. SNN makes this capability native to all accounts.
- **SNN Innovation:** Bringing full, native account abstraction to a scalable L1 DAG, making advanced account features the default, not the exception.

## 4.2. Predictable & Low Fees

Achieving low and predictable transaction fees through DAG efficiency and an optimized fee market.

- **DAG Efficiency Impact:**
  - **High Throughput:** The Synaptic Ledger's parallelism (2.1) allows many TUs to be processed concurrently. This high capacity reduces congestion, which is a primary driver of high fees on serial blockchains.
  - **Low Latency Confirmation:** Avalanche consensus provides fast (probabilistic) finality. Quicker settlement means less need for users to overbid on fees to ensure inclusion.
  - **Reduced Wasted Effort:** Unlike PoW, PoS DAGs don't have massive energy expenditure or orphaned block races that indirectly add to fee pressure.
- **Optimized Fee Market (Multi-dimensional Considerations, Algorithms):**
  - **Concept:** SNN aims for a fee market that is more nuanced than simple highest-bidder-wins, reflecting the multi-dimensional nature of resource consumption (computation, storage, bandwidth).
  - **Algorithm (Base Fee + Priority Fee - EIP-1559 like):**
    - **Base NEURON Fee:** A protocol-defined base fee per unit of computation ("gas") and per byte of TU data. This base fee can adjust dynamically based on network load (e.g., if DAG utilization is above X% for Y epochs, base fee increases; if below, it decreases). A portion of the base fee can be burned, and another portion can go to the SNN Treasury or validators.
    - **Priority NEURON Fee (Tip):** Users can add a tip to incentivize validators to include their transaction in a TU faster, especially during transient congestion. Tips go directly to the TU-issuing validator.
  - **Algorithm (Storage Fee):** Separate fee for persistent state storage (see State Rent 2.3.1). This is not part of the per-transaction execution fee but is crucial for long-term predictability.

- **Algorithm (Bandwidth Fee Component):** The base fee for TU data can implicitly cover bandwidth. For very large TUs or those requiring extensive gossiping (e.g., cross-DSC messages), a more explicit bandwidth component could be added.
- **Data Structures:**
  - `NetworkCongestionOracle (Protocol State)` : Tracks recent DAG utilization (e.g., TU count, total gas consumed, TU bytes) to adjust the base fee.
  - Each TU/UserOperation specifies `max_fee_per_gas` and `max_priority_fee_per_gas` .
- **Component Breakdown:**
  - **Fee Calculation Logic (in client SDKs and validator software).**
  - **Base Fee Adjustment Mechanism (protocol-level).**
  - **Validator TU Selection Logic (considers priority fees).**
- **Implementation Challenges & Trade-offs:**
  - **Challenge (Accurate Resource Metering):** Precisely measuring all relevant resource consumption (CPU, I/O, bandwidth, storage duration) is hard. Gas is an abstraction.
  - **Challenge (Dynamic Base Fee Stability):** Tuning the EIP-1559-like mechanism to be responsive but not overly volatile.
  - **Trade-off (Predictability vs. Market Efficiency):** Fixed fees are predictable but not market-efficient. Pure auction is efficient but unpredictable. EIP-1559 aims for a balance.
- **Security Considerations:**
  - **Validator Collusion (Fee Manipulation):** E.g., validators only including high-tip transactions even if base fee is met. Mitigated by PoS incentives and user choice of validators.
  - **Economic Exploits of Fee Market:** Design must be robust against attempts to artificially manipulate the base fee.
- **Parameterization & Configuration:**
  - Target DAG utilization for base fee.
  - Base fee adjustment speed/magnitude.
  - Initial base fee values.
  - Portion of base fee burned/treasured.
- **Integration Points:**
  - **Transaction Processing (1.3).**
  - **Execution Environment (5.1, for gas metering).**
  - **State Rent (2.3.1).**
  - **Account Abstraction (4.1, for fee payment options).**



- **Research & Prototyping Needs:**
  - Simulation of the EIP-1559-like mechanism on a DAG.
  - Developing more sophisticated multi-dimensional fee models if needed.
- **Comparison with Existing Implementations:**
  - **Ethereum (EIP-1559):** SNN adopts similar principles for base fee + priority fee.
  - **Solana, Avalanche (Low Fees due to High TPS):** SNN aims for similar low fees via its own high-throughput DAG architecture.
  - **Near (Predictable Gas Costs):** Gas units are tied to TFLOPS, aiming for more stable cost prediction. SNN can learn from this for its gas model.
  - **SNN Innovation:** Combining DAG efficiencies with an EIP-1559 style market and native AA-enabled fee delegation to provide a comprehensive low-fee, predictable environment.

#### 4.2.3. Fee Delegation/Sponsorship Mechanisms

Allowing users to have their transaction fees paid by a third party (Paymaster).

- **Concept:** As part of Account Abstraction (4.1), a user's AC can delegate fee payment to another contract, known as a "Paymaster." The Paymaster agrees to pay the NEURON fees for the UserOperation, potentially in exchange for the user paying the Paymaster in another token, or for free (sponsored transaction).
- **Specific Algorithms & Data Structures:**
  - **Data Structure (UserOperation):** Includes optional fields: `paymaster` (address of Paymaster contract), `paymasterData` (data for Paymaster to validate/execute).
  - **Interface (IPaymaster):**
    - `validatePaymasterUserOp(userOp, requiredNeuronFee)` : Called by EntryPoint/protocol. Paymaster inspects `userOp` and `paymasterData` . If it agrees to pay, it ensures it has enough NEURON to cover `requiredNeuronFee` (e.g., by having NEURON deposited with the EntryPoint or by allowing the EntryPoint to withdraw from it) and returns context for `postOp` .
    - `postOp(mode, context, actualNeuronFee)` : Called after UserOperation execution. Paymaster can execute logic, e.g., charge the user in an alternative token. `mode` indicates if op was success/failure.
  - **Algorithm (Transaction Flow with Paymaster):**
    1. UserOperation includes Paymaster details.
    2. EntryPoint/Protocol first calls `AC.validateTransaction()` .
    3. If AC validation succeeds, EntryPoint/Protocol calls `Paymaster.validatePaymasterUserOp()` .

4. If Paymaster validation succeeds, EntryPoint/Protocol executes the main UserOperation.
  5. EntryPoint/Protocol calls `Paymaster.postOp()`.
  6. Fees (in NEURON) are paid by the Paymaster to the validator/protocol.
- **Component Breakdown:**
    - **Paymaster Contract Standard/Templates.**
    - **EntryPoint/Protocol logic to handle Paymaster interactions.**
    - **Wallet SDK support for discovering and using Paymasters.**
  - **Implementation Challenges & Trade-offs:**
    - **Challenge (Paymaster Security):** Paymasters manage funds and complex logic; bugs can be costly.
    - **Challenge (Paymaster Discovery & Trust):** Users need to find and trust Paymasters.
    - **Trade-off (Complexity vs. Flexibility):** The Paymaster flow adds complexity but offers great UX flexibility.
  - **Security Considerations:**
    - **Paymaster Insolvency:** If a Paymaster agrees to pay but can't, the transaction fails. EntryPoint needs to manage Paymaster NEURON stakes/deposits.
    - **Griefing Paymasters:** Users spamming Paymasters with ops they won't actually pay for in alt-tokens. `validatePaymasterUserOp` must be robust.
  - **Integration Points:** Account Abstraction (4.1), Fee Market (4.2).
  - **Comparison with Existing Implementations:** EIP-4337 (defines Paymasters for Ethereum AA). SNN implements this natively. Biconomy, ZeroDev (provide Paymaster services on Ethereum).

### 4.3. Human-Readable Addressing: SNN Name Service (SNS)

Implementing a decentralized naming system (like ENS) for SNN addresses.

- **Concept:** A system that allows users to map human-readable names (e.g., `alice.snn`) to SNN account addresses (AC addresses), and potentially other data like avatar URLs, social media handles, etc.
- **Specific Algorithms & Data Structures:**
  - **Data Structure (SNS Registry Contract):**
    - Stores mappings: `hash(name) -> { owner_AC, resolver_contract_address, TTL }`.
    - Manages name registration, renewal, transfer.
    - Uses a hierarchical naming system (e.g., `.snn` as top-level, users register second-level domains).

- **Data Structure (Resolver Contract):**
  - Separate contract per name (or a shared one). Owner of the name configures it.
  - Stores actual records: `hash(name), record_type (e.g., "ADDRESS", "CONTENT_HASH", "TEXT") -> value`.
  - Implements a standard `resolve(name, record_type)` function.
- **Algorithm (Name Registration):**
  - Auction mechanism (for popular/short names) or fixed-price for longer names.
  - Commit-reveal scheme to prevent front-running in auctions.
  - Registrations are for a limited period, require renewal (pays fees to SNS DAO/Treasury).
- **Algorithm (Name Resolution):**
  1. Client wants to resolve `alice.snn`.
  2. Client queries SNS Registry for `alice.snn` -> gets resolver contract address.
  3. Client queries resolver contract for `alice.snn` with type "ADDRESS" -> gets SNN AC address.
- **Component Breakdown:**
  - **SNS Registry Contract (SNN Smart Contract).**
  - **Public Resolver Contract Template (SNN Smart Contract).**
  - **SNS DAO/Governance Contracts (for managing .snn TLD, fees, policies).**
  - **Integration into SNN Wallets and dApps (to display names, allow sending to names).**
  - **Off-chain name search/discovery services.**
- **Implementation Challenges & Trade-offs:**
  - **Challenge (Name Squatting):** Designing fair registration mechanisms.
  - **Challenge (Scalability of Registry/Resolver):** Efficient storage and querying.
  - **Trade-off (Centralization of TLD vs. User Freedom):** `.snn` TLD managed by SNS DAO. Users can create their own subdomains.
- **Security Considerations:**
  - **Bugs in Registry/Resolver Contracts.**
  - **Governance Capture of SNS DAO (controlling TLD).**
  - **Phishing with similar-looking names (Unicode normalization, UI warnings).**
- **Parameterization & Configuration:**
  - Name registration fees, renewal fees, auction parameters.

- Supported record types.
- **Integration Points:** Wallets, dApps, SNN Account Model.
- **Comparison with Existing Implementations:** ENS (Ethereum Name Service), Unstoppable Domains, Handshake. SNN aims for an ENS-like system, tightly integrated with its native AA.

#### 4.4. Simplified Onboarding

Technical aspects of enabling social logins, seedless recovery, and user-friendly wallet standards.

- **Social Logins (via OpenID Connect / OAuth2 and MPC Wallets):**
  - **Concept:** Users can use existing Web2 credentials (Google, X, etc.) to create/access an SNN account. This typically involves a non-custodial MPC wallet provider.
    1. User signs in with Google to MPC Wallet Provider.
    2. Provider uses MPC to create an SNN AC for the user. One share of the AC's key is held by the provider (encrypted by user's social login), another by user's device, maybe a third for recovery.
    3. User operations are authorized by interaction between device share and provider share (re-authenticated via social login).
  - **SNN Role:** SNN itself doesn't handle social logins. It supports ACs whose `validateTransaction` can verify signatures from MPC-derived keys. SNN fosters an ecosystem of such wallet providers.
  - **Security:** Relies on security of social login provider AND MPC wallet provider. User education is key.
- **Seedless Recovery (Leveraging AA):**
  - **Social Recovery (4.1):** ACs can configure guardians for recovery. No seed phrase needed if guardians are set up.
  - **MPC-based Recovery:** MPC wallet providers can offer recovery mechanisms tied to social logins or other user-verifiable factors, to recover their share of the key.
  - **Device-based Security:** Using device biometrics (fingerprint, FaceID) + secure enclaves to protect a key share, with other shares recoverable via social means.
- **Wallet Standards:**
  - **SNN WalletConnect Standard:** Adapting WalletConnect protocol for SNN's AA model (UserOperations instead of simple transactions).
  - **Standardized JSON-RPC API for ACs:** For dApps to query AC properties (signers, supported methods, etc.).
  - **Promoting CAIPs (as in 3.3).**

- **Implementation Challenges:**
    - Building trust in MPC wallet providers for social login.
    - Standardizing AA-aware wallet interactions.
  - **SNN's Role:** Provide the flexible AA foundation (4.1) that enables these UX improvements. Foster a competitive ecosystem of wallets and MPC providers. Support standardization efforts.
  - **Comparison:** Web3Auth, Torus Wallet (MPC social logins). Argent (social recovery). SNN enables these at the protocol level for any wallet to build upon.
- 

## 5. Developer Ecosystem & Smart Contracts

Facilitating a rich developer ecosystem through a robust smart contract platform, language choices, and comprehensive tooling.

### 5.1. Smart Contract Platform

Details of the execution environment and state model for SNN smart contracts.

#### 5.1.1. Execution Environment: WebAssembly (WASM)

- **VM Selection/Customization:**
  - **Choice:** Wasmer, Wasmtime, or a similar high-performance, secure WASM runtime. The choice will be based on performance benchmarks, security features, gas metering capabilities, and ease of integration.
  - **Customization:**
    - **Host Functions:** SNN protocol will expose a set of host functions callable from WASM smart contracts. These provide access to SNN-specific features:
      - Reading/writing contract state (SMT access).
      - Getting TU context (timestamp, issuer, current AC).
      - Calling other contracts (ACs or utility contracts).
      - Emitting events.
      - Accessing cryptographic precompiles (BLAKE3, EdDSA verify, ECDSA verify, etc.).
      - Interacting with Fee Market / Paymasters.
      - Accessing SNN Name Service.
      - Initiating cross-DSC/cross-chain messages.
    - **Gas Metering:** Inject gas counters into WASM bytecode during compilation or use runtime gas metering provided by the WASM

engine. Every WASM instruction and host function call will have a defined gas cost.

- **Deterministic Execution:** Ensure the WASM runtime and host functions are strictly deterministic. Disable features like floating-point operations (or use deterministic soft-float libraries), threads, and direct system calls.
- **Gas Model:**
  - **Gas Units:** Abstract units representing computational cost.
  - **Gas Schedule:** A table mapping each WASM opcode and SNN host function to its gas cost. This schedule is a critical network parameter, updatable by SNN governance.
  - **Transaction Gas Limit:** Each UserOperation/transaction specifies a gas limit. Execution stops if limit exceeded.
  - **Block/TU Gas Limit (if applicable):** The Synaptic Ledger might not have strict per-TU gas limits like Ethereum blocks, but validators will naturally be limited by processing time. DSCs might implement local TU/epoch gas limits.
- **Component Breakdown:**
  - **WASM Runtime Integration Layer.**
  - **SNN Host Function Implementation.**
  - **Gas Metering Injector/Module.**
  - **Smart Contract Deployment & Update Logic.**
- **Implementation Challenges:**
  - **Secure WASM Runtime:** Ensuring sandbox security and preventing exploits.
  - **Deterministic Host Functions:** Careful design and testing.
  - **Accurate Gas Scheduling:** Balancing cost fairness with preventing DoS.
  - **WASM Contract Size Limits:** Managing on-chain storage of large contracts.
- **Security Considerations:**
  - **WASM Vulnerabilities (e.g., JIT bombs, sandbox escapes):** Rely on well-audited runtimes.
  - **Reentrancy Attacks:** Follow best practices (checks-effects-interactions). SNN's state model might offer some reentrancy protection (see below).
  - **Infinite Loops:** Gas limit handles this.
- **Comparison:** Ethereum (EVM -> EWASM research), Polkadot (WASM), NEAR (WASM), Solana (BPF, moving to WASM), Dfinity (WASM). WASM is the industry trend for high-performance, language-agnostic smart contracts. SNN aligns with this.

### 5.1.2. State Model for DAGs (Actor Model, STM exploration, CRDTs)

How smart contract state is managed and interacts with the WASM execution environment, especially in a DAG context.

- **Primary Model: Actor Model + SMT per Contract/Account:**
  - **Concept:** Each smart contract (including Account Contracts) is an “actor” with its own private state.
    - **State Storage:** Each contract’s state is stored in its own dedicated Sparse Merkle Tree (SMT). The root of this SMT is part of the contract’s account data in the global SNN state SMT.
    - **Message Passing:** Contracts interact by sending asynchronous messages (internal transactions/calls). A call from Contract A to Contract B is a message.
    - **Serial Execution within an Actor:** Messages processed by a single actor (contract) are executed serially against its state, preventing internal race conditions for that actor.
  - **Interaction with WASM:** WASM host functions allow a contract to read/write key-value pairs within its *own* SMT. When Contract A calls Contract B, the SNN execution environment handles the context switch, loads Contract B’s SMT (if needed), and executes B’s WASM code.
- **Concurrency & DAGs:**
  - The DAG (Synaptic Ledger) provides a partially ordered set of TUs. Transactions within a TU are typically ordered. Consensus (Avalanche) finalizes this order.
  - **Parallel Execution Potential:** If two transactions in different TUs (or even same TU if independent) affect different actors (contracts), their execution can be parallelized by the SNN node after consensus, as they operate on distinct SMTs.
  - **Account-Synapses (2.1):** Reinforce this. Operations on different account-synapses (which are ACs, hence actors) can often proceed in parallel.
- **Software Transactional Memory (STM) - Exploration for Complex Interactions:**
  - **Concept:** For operations that need to atomically update state across *multiple* actors (contracts), STM could be explored. An STM system would allow a sequence of operations across contracts to be bundled into a “transaction.” If any part fails or a conflict is detected (e.g., another transaction modified intermediate state), the entire STM transaction is rolled back.
  - **SNN Use:** Could be a higher-level primitive offered to developers for specific use cases, or an internal mechanism for handling complex multi-contract calls. This is an advanced feature, not for initial launch.

- **Challenges:** Complexity, performance overhead of STM.
- **Conflict-Free Replicated Data Types (CRDTs) - Exploration for Specific State Types:**
  - **Concept:** For certain types of shared state (e.g., counters, sets, shared documents) that are frequently updated concurrently, CRDTs offer a way to ensure eventual consistency without locks or complex coordination. Each replica can update its local copy; CRDT merge functions ensure all replicas converge to the same state.
  - **SNN Use:** Could be offered as a library or built-in data types for SNN smart contracts, especially useful in highly concurrent DSCs or for off-chain + on-chain data synchronization.
  - **Challenges:** Limited to data types with available CRDT implementations. Developer understanding.
- **Component Breakdown:**
  - **SMT Management Library (for per-contract state).**
  - **Execution Dispatcher (handles inter-contract calls, context switching).**
  - **(Future) STM Engine / CRDT Library for WASM contracts.**
- **Implementation Challenges:**
  - Efficiently managing a large number of per-contract SMTs.
  - Optimizing inter-contract call overhead.
  - Designing clear semantics for developers regarding state access and concurrency.
- **Security Considerations:**
  - **Reentrancy:** While actor model helps, if Contract A calls B, and B calls A back *before* A's first call finishes updating its state, reentrancy can still occur if not handled carefully by A. SNN might enforce a “no reentrancy during SMT update” rule or provide clear developer guidance.
  - **Read-Write Conflicts (if not pure actor):** If contracts could directly read other contracts' SMTs (not just via calls), managing consistency is harder. Pure actor model (only private state writes, reads via calls or committed state) is safer.
- **Comparison:**
  - **Ethereum (Shared Key-Value Store):** All contracts share one massive state trie. Simpler for some things, harder for parallelism.
  - **Polkadot (Actor-like Contracts):** Contracts have their own storage.
  - **Dfinity (Canister Actors):** Pure actor model. SNN is similar.
  - **Hyperledger Fabric (World State per Channel):** More permissioned, but concept of isolated state.
  - **SNN Innovation:** Combining the actor model with per-contract SMTs on a DAG, designed for high parallelism and clear state isolation, while



exploring STM/CRDTs for advanced use cases.

## 5.2. Choice of Programming Languages

Supporting popular languages and potentially a domain-specific language.

- **Primary Support: Rust & Go (via WASM Compilation):**
  - **Rust:** Strong ecosystem for WASM (e.g., `wasm-pack`, `cargo-contract` for Substrate-like environments). Excellent for systems programming, safety, performance. SNN will provide first-class SDKs, libraries, and examples for writing SNN smart contracts in Rust.
  - **Go:** Growing support for WASM compilation (TinyGo). Popular for backend/systems. SNN will provide SDKs for Go.
  - **Toolchain Development/Integration:**
    - SNN-specific contract development kits (SDKs) for Rust and Go. These provide:
      - Macros/libraries for defining contract structure, storage, entry points.
      - Bindings to SNN host functions.
      - Serialization/deserialization helpers (e.g., using Borsh or similar).
      - Testing frameworks that mock the SNN environment.
    - Easy compilation to SNN-compatible WASM (e.g., scripts, IDE plugins).
- **Other Languages (via WASM):**
  - **AssemblyScript (TypeScript-like):** Compiles to WASM, popular for smart contracts (e.g., NEAR). SNN will support it and provide examples.
  - **C/C++:** Can compile to WASM (e.g., via Emscripten). Support for experienced developers.
  - **Swift, Kotlin (Experimental):** As WASM support for these matures.
  - **WASM Compilation Specifics:** SNN will document required WASM features, binary size recommendations, and best practices for compiling from various languages to ensure compatibility and efficiency.
- **“SynapseScript” (Exploratory DSL):**
  - **Concept:** A high-level, domain-specific language designed specifically for SNN smart contracts. Could offer:
    - Simplified syntax for common SNN patterns (AA, fee delegation, SMT interaction, actor calls).
    - Built-in safety features tailored to SNN’s state model.
    - Easier formal verification.
  - **Details if Pursued:**
    - Would likely compile to SNN WASM.

- Development would require significant R&D (compiler, tooling).
- Initial focus is on robust Rust/Go support. SynapseScript is a longer-term research goal. If pursued, it would be introduced alongside existing language support, not as a replacement.
- **Component Breakdown:**
  - **SNN SDKs for Rust, Go, AssemblyScript.**
  - **Documentation & Tutorials for each supported language.**
  - **(Future) SynapseScript Compiler & Toolchain.**
- **Implementation Challenges:**
  - Maintaining high-quality SDKs for multiple languages.
  - Ensuring consistent security and performance from different language toolchains targeting WASM.
  - Huge effort if SynapseScript is developed.
- **Community Engagement:** Actively support community efforts to bring more languages to SNN WASM.

### 5.3. Tooling

Providing a comprehensive suite of tools for developers.

- **SDKs (Software Development Kits):**
  - **Client SDKs (JavaScript/TypeScript, Python, Go, Rust):** For building dApps and backend services that interact with SNN (querying state, submitting UserOperations, interacting with ACs and contracts).
    - Include helpers for UserOperation creation, signing, AC deployment.
    - Abstraction for SNS resolution, fee estimation, Paymaster interaction.
  - **Contract SDKs (Rust, Go, AssemblyScript):** As described in 5.2.
- **APIs:**
  - **JSON-RPC API:** Standardized API for SNN nodes, similar to Ethereum's but adapted for SNN's DAG structure and AA model.
    - Endpoints for: sending UserOperations, querying TU status, getting AC/contract state (via SMT proofs), reading SNS, estimating fees, etc.
  - **Streaming APIs (WebSockets/gRPC):** For real-time updates on TUs, events, state changes.
- **IDE Integrations:**
  - Plugins for popular IDEs (VS Code, IntelliJ IDEA) for supported contract languages.
  - Features: syntax highlighting, code completion (using SDKs), compilation to WASM, deployment to testnet/localnet, debugging support.

- **Debuggers:**
  - **WASM Debugging:** Leverage existing WASM debugging tools, potentially with SNN-specific extensions to inspect contract state (SMTs) and execution context during local testing.
  - **Transaction Tracing:** Tools to trace execution flow of a UserOperation, including inter-contract calls and state changes, on testnets and mainnet (for developers to understand failures or performance).
- **Formal Verification Tool Integration:**
  - Encourage and support the use of formal verification tools for SNN smart contracts, especially for critical infrastructure (bridges, DAO, core AC templates).
  - This might involve:
    - Writing SNN host function specifications in a language understood by FV tools.
    - Developing contract SDK features that output models suitable for FV.
    - Partnering with FV firms/projects.
- **Local Testnet Setup:**
  - Simple CLI tool or Docker Compose setup for developers to spin up a local single-node or multi-node SNN instance for development and testing.
  - Pre-configured with accounts, deployed utility contracts (SNS, EntryPoint if used).
- **Component Breakdown:**
  - Separate teams/efforts for SDKs, API dev, IDE plugins, etc.
  - Strong documentation site.
- **Implementation Challenges:**
  - Building and maintaining a wide range of high-quality tools is a continuous effort.
  - Keeping tools synchronized with protocol upgrades.
- **Community & Ecosystem:** Foster a marketplace for third-party developer tools and services.

## 5.4. Modularity & Composability

Designing SNN and its smart contract platform for easy creation and interaction of modular components.

- **Implementation of Feature Modules (Protocol Level):**
  - SNN itself is modular (Consensus, AA, Staking, Governance, Bridges are distinct modules).
  - Smart contracts can be designed as reusable libraries/modules.

- **Standardized Interfaces (Contract Level):**
    - Promote use of clear, well-documented interfaces for SNN smart contracts (e.g., SNN's versions of ERC-20, ERC-721, ERC-1155 for tokens).
    - Standard interfaces for ACs (4.1), Paymasters (4.2.3), SNS Resolvers (4.3).
    - This allows contracts to interact predictably.
  - **Cross-Contract Call Mechanisms:**
    - SNN WASM host functions provide secure, gas-metered mechanisms for synchronous ( `call` ) and potentially asynchronous ( `send_message` for actor-style interactions) calls between contracts.
    - Clear semantics for call depth limits, error handling, return data.
    - Calls between contracts on different DSCs are asynchronous and go via the cross-DSC messaging protocol (2.2.3).
  - **Event System:**
    - Standardized event logging via a host function. Events are stored on-chain (or their hashes committed to) and can be indexed by off-chain services for dApps to query.
  - **Upgradability Patterns for Smart Contracts:**
    - Support for proxy patterns (e.g., Universal Upgradeable Proxy Standard - UUPS) for contract upgradability, managed by contract owners or DAOs.
    - SNN governance might also provide a mechanism for fixing critical bugs in widely used standard contract libraries, with opt-in from users.
  - **Benefits:** Encourages code reuse, reduces development effort, fosters a rich ecosystem of interoperable dApps.
  - **Challenges:** Ensuring interface compatibility over time, managing risks of complex interactions.
- 

## 6. Economic Model & Governance

Defining the utility of the NEURON token and the framework for SNN governance.

### 6.1. Native Utility Token (Symbol: NEURON)

Detailing the uses of NEURON within the SNN ecosystem.

#### 6.1.1. Staking & Network Security

- **Validator Contracts & Delegation Logic:**

- **Validator Staking:** To become an SNN validator (issuing TUs, participating in Avalanche consensus), entities must stake a minimum amount of NEURON in a `StakingContract`. This stake acts as collateral, slashable for misbehavior.
- **Delegation:** NEURON holders who don't want to run a validator node can delegate their NEURON to an existing validator.
  - `StakingContract` manages delegation: delegators choose a validator, lock NEURON with them.
  - Validator's total stake (own + delegated) determines their weight in consensus (probability of being sampled in Avalanche) and potentially their capacity to issue TUs or earn rewards.
  - Rewards (from transaction fees, inflation if any) earned by the validator are shared proportionally with their delegators, after the validator takes a commission.
  - Delegators share risk: if validator is slashed, a portion of delegated stake is also slashed.
- **Slashing Implementation:**
  - `SlashingContract` (or module within `StakingContract`).
  - **Slashable Offenses:**
    1. **Liveness Failures:** Prolonged downtime / non-participation in consensus (e.g., failing to respond to Avalanche queries consistently).
    2. **Safety Failures (Severe):** Verifiably signing conflicting TUs (double-spending attempts if not resolved by consensus), issuing provably invalid TUs, or confirmed malicious actions in bridge MPC/light client roles.
  - **Proof Submission:** Evidence of misbehavior (e.g., signed conflicting messages, proof of downtime from monitoring nodes) is submitted to the `SlashingContract`.
  - **Slashing Penalty:** A percentage of the validator's (and their delegators') stake is burned or sent to the SNN Treasury. The validator might also be "jailed" (temporarily or permanently removed from the active set).
  - **Attributability:** Designing robust mechanisms to correctly attribute blame, especially for liveness failures, is crucial.
- **Component Breakdown:** `StakingContract`, `SlashingContract`, validator client software (for participation), delegator wallet interfaces.
- **Security:** Slashing is the primary economic disincentive against attacks. Minimum stake and unbonding periods prevent rapid withdrawal after misbehavior.

### 6.1.2. Paying Transaction Fees

- **Collection, Burning, Distribution Mechanisms:**
  - **Fee Currency:** Primary fee token is NEURON. Account Abstraction (4.1, 4.2.3) allows users to pay with other tokens via Paymasters, but Paymasters ultimately settle fees in NEURON with the protocol/validators.
  - **Base Fee (EIP-1559 style, from 4.2):**
    - Collected by the protocol.
    - **Burning:** A significant portion of the base fee (e.g., 50-100%) is burned, making NEURON deflationary with network usage.
    - **Distribution (Optional):** The remaining portion (if not 100% burned) could go to:
      - SNN Treasury (for ecosystem development, governed by NEURON holders).
      - Validators (as a baseline reward, supplementing tips). This needs careful balancing with inflation if any.
  - **Priority Fee (Tip, from 4.2):** Paid directly to the validator that includes the UserOperation/transaction in their TU.
  - **Storage Fees (State Rent, from 2.3.1):** Collected in NEURON. Can be burned or go to Treasury/validators responsible for state maintenance.
- **Component Breakdown:** Fee handling logic in protocol, Paymaster contracts, Treasury contract.

### 6.1.3. Governance Participation

- **Token Voting Contracts for On-Chain Proposals:**
  - NEURON holders can vote on SNN governance proposals.
  - Voting power is proportional to the amount of NEURON staked or locked for voting (details in 6.2).
  - Proposals can cover: protocol parameter changes (e.g., fee constants, slash penalties), software upgrades, Treasury fund allocation, SNN Council elections, etc.
  - `GovernanceContract` manages proposal submission, voting periods, tallying, and execution of passed proposals.
- **Component Breakdown:** `GovernanceContract` , off-chain proposal discussion forums.

### 6.1.4. Powering Specific Network Services

- **DSC Registration/Bonding:** Projects or entities wishing to create a new Dynamic Synaptic Cluster (DSC, see 2.2) might need to pay a registration fee in NEURON or bond a certain amount of NEURON, which could be

returned if the DSC is decommissioned gracefully or slashed if it misbehaves.

- **Premium SNN Name Service (SNS) Names:** While standard SNS names (4.3) might have regular fees, very short, desirable, or special TLDs (if allowed beyond `.snn`) could be auctioned for NEURON, with proceeds going to the SNS DAO or SNN Treasury.
- **Oracle Services (SNN-Native):** If SNN fosters native oracle networks (3.2.1), NEURON would likely be used for staking by oracle node operators and for paying for oracle data requests.
- **Other Future Services:** Any new protocol-level service (e.g., specialized storage networks, computation networks built on SNN) could use NEURON for access fees, staking, or quality-of-service bonds.
- **Component Breakdown:** Specific contracts/modules for each service that integrates NEURON.

## 6.2. Governance Framework

Defining the mechanisms for decision-making and protocol evolution. SNN aims for a robust, decentralized on-chain governance model.

### 6.2.1. On-Chain Binding Votes

- **Proposal Submission:**
  - Any NEURON holder (or one holding above a certain threshold of NEURON) can submit a governance proposal via the `GovernanceContract`.
  - Proposals must typically be accompanied by a deposit of NEURON, which is returned if the proposal meets minimum criteria (e.g., not spam, reaches a quorum for voting) or slashed if it's deemed malicious or frivolous.
  - Proposals should include: description, motivation, and executable code (if it's a parameter change or contract upgrade) or clear instructions.
- **NEURON-Weighted Voting Process:**
  - **Voting Period:** A fixed duration during which NEURON holders can cast their votes (Yes, No, Abstain).
  - **Eligibility:** Typically, NEURON that is staked (for validation or delegation) or explicitly locked in the `GovernanceContract` for voting is eligible. This prevents vote selling with liquid tokens.
  - **Quorum:** A minimum percentage of eligible NEURON must participate in the vote for the proposal to be considered valid.
  - **Threshold:** A minimum percentage of "Yes" votes (e.g., >50% of non-abstain votes, or >66% for critical changes) is required for a proposal to pass.



- **Execution of Passed Proposals:**

- If a proposal passes (quorum and threshold met), the `GovernanceContract` can autonomously execute it after a “timelock” period.
- **Timelock:** A delay between when a vote passes and when its changes are enacted. This gives the community time to react, prepare for changes, or even trigger emergency measures if a malicious proposal somehow passes (e.g., via SNN Council veto, see below).
- **Executable Payloads:** For parameter changes or smart contract upgrades, the proposal includes the actual code/parameters. The `GovernanceContract` has administrative privileges (e.g., `set_fee_parameter()` , `upgrade_contract_to(new_code_hash)` ) over relevant SNN system contracts, callable only by a passed proposal.

- **Component Breakdown:** `GovernanceContract` (handles proposals, voting, timelock, execution), `TreasuryContract` (managed by governance).

### 6.2.2. SNN Council (Optional – “Security Council” / “Technical Steering”)

A small, elected body with limited, clearly defined powers, primarily for safety or rapid response, or technical guidance. This is a common pattern to balance pure token voting with expert oversight.

- **Election Mechanisms:**

- Elected by NEURON holders via the `GovernanceContract` .
- Candidates stake NEURON to run.
- Regular election cycles (e.g., every 6-12 months).
- Could use approval voting or ranked-choice voting.

- **Powers (Strictly Limited):**

- **Emergency Veto:** The Council might have the power to veto a passed governance proposal during its timelock period if it’s deemed critically harmful (e.g., a bug in the proposal code). This veto itself could be overridden by a supermajority NEURON vote.
- **Fast-Tracking Proposals:** For urgent bug fixes or critical upgrades, the Council might be able to propose changes with a shorter voting/timelock period, but still requiring NEURON holder approval.
- **Technical Advisory:** Provide non-binding recommendations on proposals.
- **NOT for unilateral changes or Treasury control.**

- **On-Chain Representation:** Council members’ AC addresses are stored in the `GovernanceContract` . Multi-sig actions by the Council (e.g., a veto) are verified on-chain.



- **Trade-off:** Introduces a point of (elected) centralization, but can protect against unforeseen issues with pure token voting or enable faster responses. Transparency is key.

### 6.2.3. Technical Committee (Advisory)

- **Role:** A body of core developers, researchers, and security experts recognized by the SNN Foundation or community.
- **Functions (Primarily Off-Chain or Advisory On-Chain):**
  - Review technical feasibility and security implications of proposals.
  - Publish analyses and recommendations.
  - Maintain the SNN software repository.
  - Propose emergency fixes (to be approved by Council/Governance).
  - They do *not* have direct on-chain execution power over the protocol beyond what any NEURON holder has, unless also elected to the SNN Council or via Foundation multi-sigs for initial deployments/critical patches.
- **Selection:** Appointed by SNN Foundation, or through community reputation, not typically via on-chain election for this purely advisory role.

### 6.2.4. Ecosystem Fund/Treasury

- **Smart Contract for Treasury Management:**
    - `TreasuryContract` holds NEURON (and potentially other tokens).
    - Funded by:
      - A portion of NEURON base fees (if not all burned).
      - A portion of inflation (if SNN has inflation for staking rewards).
      - Proceeds from SNS name auctions, DSC registration fees.
      - Donations.
  - **Governed by NEURON Holders:**
    - Allocations from the Treasury (e.g., for grants, bounties, core development, marketing, liquidity incentives) are decided by NEURON holder governance proposals submitted to and voted on via the `GovernanceContract`.
    - The `TreasuryContract` only releases funds upon successful execution of such a proposal.
  - **Transparency:** All Treasury balances and transactions are publicly visible on-chain.
-

## 7. Learning & Differentiation / SNN's Unique Value Propositions

Highlighting how SNN's *implementation choices* differentiate it or address shortcomings of existing systems more concretely.

- **Hybrid DAG + Avalanche Consensus (Section 1.1):**
  - **Differentiation:** While others use DAGs (IOTA, Fantom, Hedera) or Avalanche consensus (Avalanche platform), SNN's specific *hybrid* design (influences from Tangle, Block-Lattice for account-synapses, Avalanche for finality) aims for a unique balance of high throughput, low latency, MEV resistance (via threshold encryption & fair ordering integrated at this layer), and robust PoS security.
  - **Addresses Shortcomings:** Overcomes the probabilistic finality concerns of early Tangle by adding Avalanche. Addresses potential ordering issues in pure DAGs with consensus timestamping. Provides a more decentralized and potentially fairer alternative to single-leader systems.
- **Native Account Abstraction (Section 4.1):**
  - **Differentiation:** Unlike Ethereum's EIP-4337 (an overlay), SNN builds AA into the core protocol. Every account *is* a smart contract. This is similar to StarkNet/zkSync but on a L1 DAG.
  - **Addresses Shortcomings:** Eliminates EOA limitations (single key, no social recovery, fee payment inflexibility) from day one for all users, leading to vastly improved UX without relying on centralized bundlers or relayers for core AA functionality.
- **Dynamic Synaptic Clusters (DSCs) with DAG Synergy (Section 2.2):**
  - **Differentiation:** SNN's sharding is envisioned as dynamic (formed by load/demand) and deeply synergistic with the main Synaptic Ledger (DAG-to-DAG anchoring). This is akin to Avalanche Subnets but with a DAG foundation for both mainnet and subnets.
  - **Addresses Shortcomings:** Offers more flexible and potentially more integrated scaling than fixed shards or L2s that have more distinct security/trust assumptions from L1. Aims for smoother composability between DSCs via the main DAG.
- **Integrated State Management (State Rent, Hierarchical Storage, Partial Statelessness) (Section 2.3):**
  - **Differentiation:** SNN plans for sustainable state management from the outset, combining state rent (like NEAR), hierarchical storage (hot/warm/cold), and a path towards partial stateless validation.
  - **Addresses Shortcomings:** Tackles state bloat, a major issue for mature L1s like Ethereum, proactively, aiming to keep validator costs low and decentralization high long-term.
- **Comprehensive MEV Resistance Strategy (Section 1.1.4):**

- **Differentiation:** Combining threshold encryption for transaction confidentiality with consensus-based fair ordering at the DAG/consensus layer.
- **Addresses Shortcomings:** Many L1s struggle with MEV or rely on off-chain solutions (Flashbots). SNN aims for base-layer MEV mitigation, promoting a fairer environment for users.
- **WASM with Actor Model & Per-Contract SMTs (Section 5.1):**
  - **Differentiation:** Provides a high-performance WASM runtime with a state model (actors, each with own SMT) that is well-suited for parallel execution on a DAG and offers strong state isolation.
  - **Addresses Shortcomings:** Avoids the shared state contention of Ethereum's EVM model, potentially leading to better scalability and fewer reentrancy concerns if designed carefully.
- **User-Centric Fee Model (Section 4.2):**
  - **Differentiation:** EIP-1559 style base+priority fees, plus native fee delegation via AA, and separate state rent.
  - **Addresses Shortcomings:** Aims for more predictable and lower fees than pure auction markets, and greater flexibility in payment than traditional systems.
- **Robust On-Chain Governance with Optional Council (Section 6.2):**
  - **Differentiation:** Clear mechanisms for binding on-chain voting for NEURON holders, coupled with an *optional, elected* Council with strictly limited powers for safety/agility.
  - **Addresses Shortcomings:** Balances pure token democracy (which can be slow or susceptible to whale dominance for nuanced decisions) with expert oversight, aiming for adaptable and secure long-term governance.

By focusing on these implementation choices, SNN intends to deliver a next-generation blockchain that is not just incrementally better but offers step-change improvements in scalability, usability, fairness, and sustainability.

---